

Case Study V: A Help-Desk Service

Prof. Daniel A. Menascé
Department of Computer Science
George Mason University
www.cs.gmu.edu/faculty/menasce.html

1

© 2004 D. A. Menascé. All Rights Reserved.

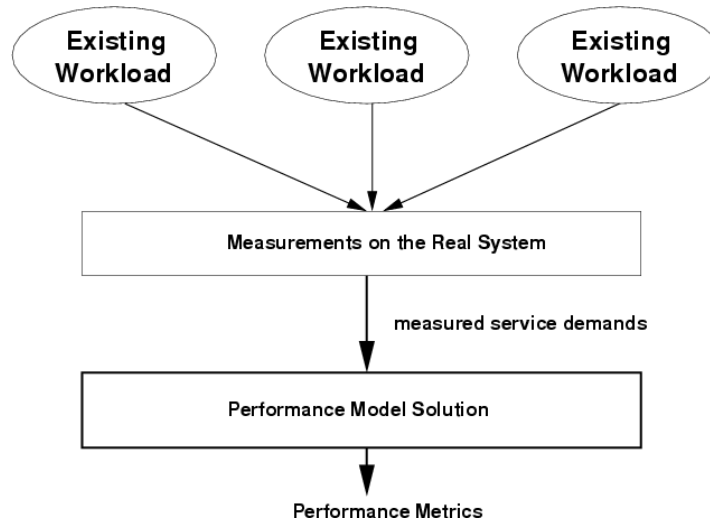
Copyright Notice

- Most of the figures in this set of slides come from the book “Performance by Design: computer capacity planning by example,” by Menascé, Almeida, and Dowdy, Prentice Hall, 2004. It is strictly forbidden to copy, post on a Web site, or distribute electronically, in part or entirely, any of the slides in this file.

2

© 2004 D. A. Menascé. All Rights Reserved.

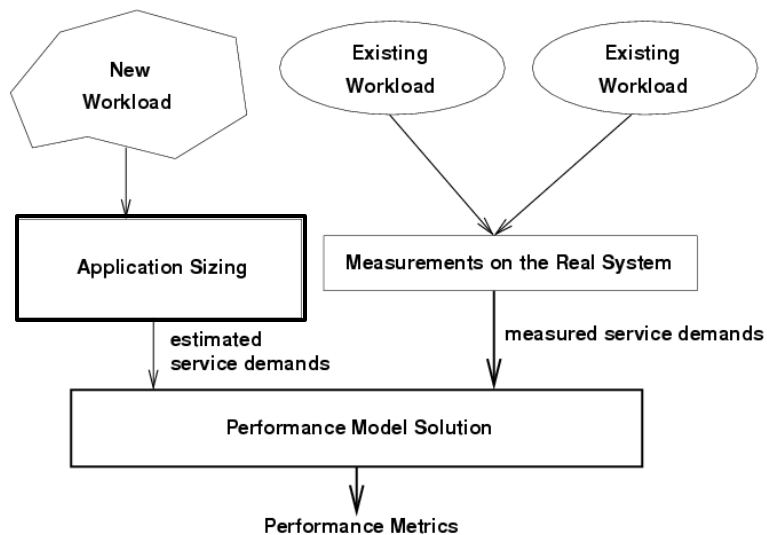
Capacity Planning for Existing Workloads



© 2004 D. A. Menascé. All Rights Reserved.

3

Capacity Planning for New Workloads



© 2004 D. A. Menascé. All Rights Reserved.

4

Software Performance Engineering (SPE)

- A set of methods and techniques to design software systems in such a way that they will adhere to performance requirements once implemented.
- SPE has to be applied at different stages during the life cycle of a software system so that estimates on services demands are continuously refined.
- Software Performance Engineers need to provide feedback to developers during the software development stage to improve the performance of the system in question.

5

© 2004 D. A. Menascé. All Rights Reserved.


The Help-Desk System

- A high-tech company has 21,600 employees with access to a variety of computing resources.
- A new help-desk application is being designed:
 - Access to a DB of Frequently Asked Questions (FAQ). Keyword-based search.
 - Creation of help tickets. If the FAQ DB does not solve the problem a help ticket is created.
 - Tracking and verification of the status of open help tickets.

6

© 2004 D. A. Menascé. All Rights Reserved.

The Help-Desk System



© 2004 D. A. Menascé. All Rights Reserved.

7

Data for Application Sizing

Number of employees	21,600
Percent employees that access FAQ DB per day	0.20
Percent access to FAQ DB at peak period	0.80
Avg. number of queries to FAQ DB per access to FAQ DB	4.00
Percent of accesses to the FAQ DB that generate ticket creation	0.35
Percent of employees that create new tickets/day without going to the FAQ database	0.15
Percent of new ticket creation at peak period	0.54
Percent of employees that inquire about new ticket status/day	0.10
Percent of status inquiries at the peak period	0.90
Peak period duration (in seconds)	7,200

Arrival rates during peak period (requests/sec)

FAQ queries	1.9200
New ticket creation	0.4110
Ticket status	0.2700

Period (in days) during which tickets are kept	365
--	-----

© 2004 D. A. Menascé. All Rights Reserved.

8

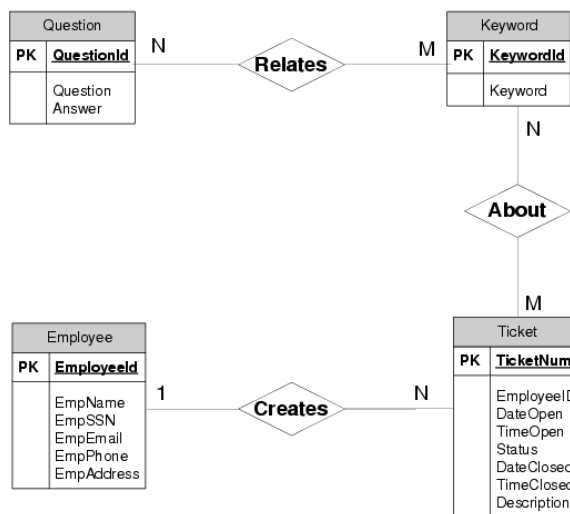
Workload Intensities

$$\begin{aligned}\lambda_{FAQ} &= \frac{N_e \times f_{FAQ} \times P_{FAQ} \times N_q}{T} \\ &= \frac{21,600 \times 0.2 \times 0.8 \times 4}{7,200} = 1.92 \text{ requests/sec} \\ \lambda_{ticket} &= \frac{N_e \times [(f_{FAQ} \times P_{FAQ} \times t_{FAQ}) + (f_t \times P_t)]}{T} \\ &= \frac{21,600 \times [(0.2 \times 0.8 \times 0.35) + (0.15 \times 0.54)]}{7,200} = 0.41 \text{ requests/sec} \\ \lambda_{status} &= \frac{N_e \times f_s \times P_s}{T} = \frac{21,600 \times 0.1 \times 0.9}{7,200} = 0.27 \text{ requests/sec} \quad (9.2.1)\end{aligned}$$

9

© 2004 D. A. Menascé. All Rights Reserved.

Entity-Relationship Model for the Database Design



10

© 2004 D. A. Menascé. All Rights Reserved.

Database Design for the Help-Desk System

Question	
PK	<u>QuestionId</u>
	Question Answer

Keyword	
PK	<u>KeywordId</u>
	Keyword

KeywordQuestion	
PK	<u>KeywordId, QuestionId</u>

Employee	
PK	<u>EmployeeId</u>
	EmpName EmpSSN EmpEmail EmpPhone EmpAddress

Ticket	
PK	<u>TicketNum</u>
	EmployeeId DateOpen TimeOpen Status DateClosed TimeClosed Description

TicketEmployee	
PK	<u>TicketId, EmployeeId</u>

TicketKeyword	
PK	<u>TicketId, KeywordId</u>

11

© 2004 D. A. Menascé. All Rights Reserved.

Database Table Data

Database Table	Cardinality	Row Size(bytes)	Total size(Kbytes)
Question	10,000	1000	10,000
Keyword	500	50	25
Employee	21,600	140	3,024
Ticket	1,734,480	260	450,965
KeywordQuestion	100,000	20	2,000
TicketEmployee	1,734,480	20	34,690
TicketKeyword	8,672,400	20	173,448

Assumptions:

- A ticket is kept in the DB for one year.
- One question is associated with 10 keywords, on average.
- One ticket has 5 keywords associated to it, on average.
- Questions are 1000 bytes long.
- Keywords are 50 bytes long.

12

© 2004 D. A. Menascé. All Rights Reserved.

Specifying Transaction Logic for SPE Purposes

- Use a language that captures the major structural components of a transaction.
 - Loops and average number of times executed.
 - Branch statements and branching probabilities.
 - Switch statements and case probabilities.
 - Database access (i.e., select and update) statements.
- Estimate number of I/Os per transaction and estimated the CPU time from the number of I/Os using benchmark data.
- Example of such a language: *Clisspe*

13

© 2004 D. A. Menascé. All Rights Reserved.

Transaction Logic for Query on FAQ Database

```
01 loop #KeywordsPerQuery
02   ! obtain keyword id for a given a keyword
03   select from Keyword where Keyword;
04   ! obtain all question ids for a given keyword id
05   select from KeywordQuestion where KeywordId;
06   ! access the Question table to retrieve selected questions
07   loop #QuestionsPerKeyword
08     select from Question where QuestionId;
09   end_loop;
10 end_loop;
```

14

© 2004 D. A. Menascé. All Rights Reserved.

Transaction Logic for the Creation of a New Ticket

```
01 ! Access the employee table to verify existence of an employee
02 select from Employee where EmployeeId;
03 if #ProbValidEmployee
04 then ! create ticket record
05     update Ticket num_rows= 1;
06     ! create a record in the TicketEmployee table
07     update TicketEmployee num_rows= 1;
08     ! create records in the TicketKeyword table
09     update TicketKeyword num_rows= #KeywordsPerTicket;
10 end_if;
```

15

© 2004 D. A. Menascé. All Rights Reserved.

Transaction Logic for Viewing the Status of Open Tickets

```
01 ! Get all ticket ids for the employee
02 select from TicketEmployee where EmployeeId;
03 ! Retrieve all open tickets for the employee
04 loop #TicketsPerEmployee
05     select from Ticket where TicketId, Status
06 end_loop;
```

16

© 2004 D. A. Menascé. All Rights Reserved.


Estimating Number of I/Os

- Estimate the number of I/Os per database access statement.
 - Consider the existence and types of indexes on the tables.
 - Estimate the number of index blocks accessed.
 - Estimate the number of data pages accessed.

17

© 2004 D. A. Menascé. All Rights Reserved.

B*-Tree Indexes



18

© 2004 D. A. Menascé. All Rights Reserved.

B-Tree Calculations

Height of the B-tree h is calculated as:

$$h = \lceil \log_k L \rceil$$

Number of entries in the B-tree L

Fanout of the B-tree k is calculated as:

$$k = \left\lfloor \frac{(\text{PageSize} - \text{HeaderSize}_i) \times f_i}{ks + rp} \right\rfloor$$

Fill factor: % of useful space of an index page used to store keys and row pointers. f_i

Size (in size) of a key in the index ks

Size (in size) of a row pointer rp

19

© 2004 D. A. Menascé. All Rights Reserved.

Indexes for the Database Tables

PageSize (bytes)	2,048
Header size (bytes)	20
Percent useful space	0.71
Size (bytes) row pointer	4

Index Info:

Table	IndexKey	KeySize(bytes)	Fanout	Height
Question	QuestionId	10	102	2
Keyword	KeywordId	10	102	2
Employee	EmployeeId	10	102	3
Ticket	TicketId	10	102	4
KeywordQuestion	(KeywordId, QuestionId)	20	59	3
KeywordQuestion	KeywordId	10	102	3
TicketEmployee	(TicketId, EmployeeId)	20	59	4
TicketEmployee	EmployeeId	10	102	4
TicketKeyword	(TicketId, KeywordId)	20	59	4

20

© 2004 D. A. Menascé. All Rights Reserved.

Estimating No. I/Os Due to a Select Statement

- Case 1: No index is available. Do a table scan (all data pages have to be accessed).

$$ndp_T = \left\lceil \frac{\text{NumRows}_T}{nrp_T} \right\rceil$$

No. rows of T that can be stored in a data page.

- Case 2: An index is available. The number of I/Os is the height of the B-tree minus 1 (the root is assumed to be in memory) + number of data pages retrieved.

21

© 2004 D. A. Menascé. All Rights Reserved.

Database Table Sizes

Data Page Calculations:

Database Table	RowSize(bytes)	Number Rows Per Page	Number Data Pages
Question	1000	1	10,000
Keyword	50	31	17
Employee	140	11	1,964
Ticket	260	6	289,080
KeywordQuestion	20	73	1,370
TicketEmployee	20	73	23,760
TicketKeyword	20	73	118,800

22

© 2004 D. A. Menascé. All Rights Reserved.

Computing Number of I/Os

KeywordsPerQuery	2
QuestionsPerKeyword	20
ProbValidEmployee	0.9
KeywordsPerTicket	5
TicketsPerEmployee	80.3

Transaction: FAQ

Statement no.	Statement type	Table	Where	Index	No. of execs	Prob. Exec	Index I/Os	Data I/Os	Total I/Os
3	select	Keyword	Keyword	none	2	1.0	0	17	34
5	select	KeywordQuestion	KeywordId	KeywordId	2	1.0	2	20	44
8	select	Question	QuestionId	QuestionId	40	1.0	1	1	80
									158

Statement 3: Table scan
 $2 \times 1.0 \times 17 = 34$

Statement 5: Index-based retrieval
 $2 \times 1.0 \times ((3-1)+20) = 44$

23

© 2004 D. A. Menascé. All Rights Reserved.

Computing Number of I/Os

KeywordsPerQuery	2
QuestionsPerKeyword	20
ProbValidEmployee	0.9
KeywordsPerTicket	5
TicketsPerEmployee	80.3

Transaction: New Ticket

Statement no.	Statement type	Table	Where	Index	No. of execs	Prob. Exec	Index I/Os	Data I/Os	Total I/Os
2	select	Employee	EmployeeId	EmployeeId	1	1.0	2	1	3.0
5	update	Ticket		TicketId	1	0.9	3	1	3.6
7	update	TicketEmployee		(TicketId,EmployeeId)	1	0.9	3	1	3.6
				EmployeeId	1	0.9	3	0	2.7
9	update	TicketKeyword		(TicketId,KeywordId)	1	0.9	15	5	18.0
									30.9

Statement 2: index-based retrieval
 $1 \times 1.0 \times ((3-1)+1) = 3$

24

© 2004 D. A. Menascé. All Rights Reserved.

Computing Number of I/Os

KeywordsPerQuery	2
QuestionsPerKeyword	20
ProbValidEmployee	0.9
KeywordsPerTicket	5
TicketsPerEmployee	80.3

Transaction: View Tickets

Statement no.	Statement type	Table	Where	Index	No. of execs	Prob. Exec	Index I/Os	Data I/Os	Total I/Os
2	select	TicketEmployee	EmployeeId	EmployeeId	1	1.0	3	80.3	83.3
5	select	Ticket	TicketId, Status	TicketId	80.3	1.0	3	1.0	321.2
									404.5

Statement 5: index-based retrieval
 $80.3 \times 1.0 \times ((4-1)+1) = 321.2$

A very large number of I/Os!

25

© 2004 D. A. Menascé. All Rights Reserved.

Computing Service Demands

Service Demand Computations:

CPU time per 2Kbyte page 0.0015 sec/IO
 Service time per I/O (sec) 0.008 sec/IO

	FAQ	Create Ticket	View Tickets
CPU	0.237	0.046	0.607
Disk	1.264	0.2472	3.236

Very large service demand.
 Disk utilization exceeds 100%.

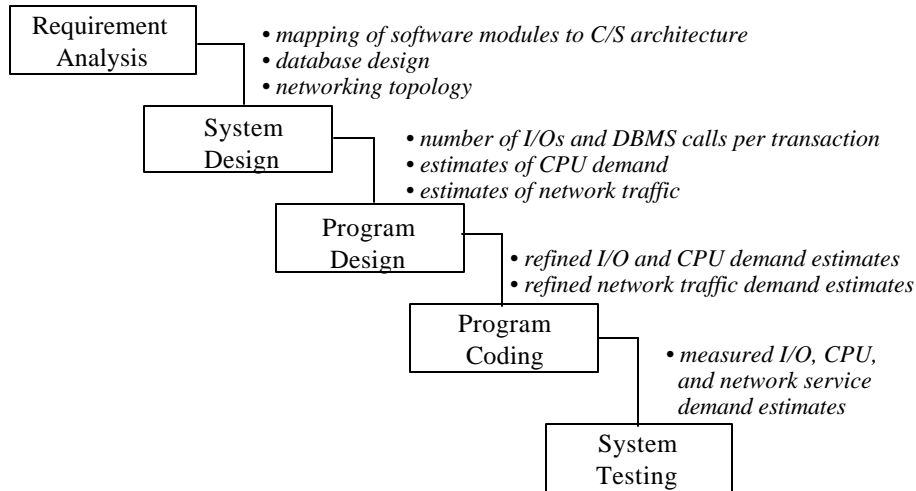
Solution: instead of storing open and closed tickets in the same table, keep only open tickets in the Tickets table and create another table to archive closed tickets for one year.

26

© 2004 D. A. Menascé. All Rights Reserved.

The Software Development Life Cycle and Inputs to SPE

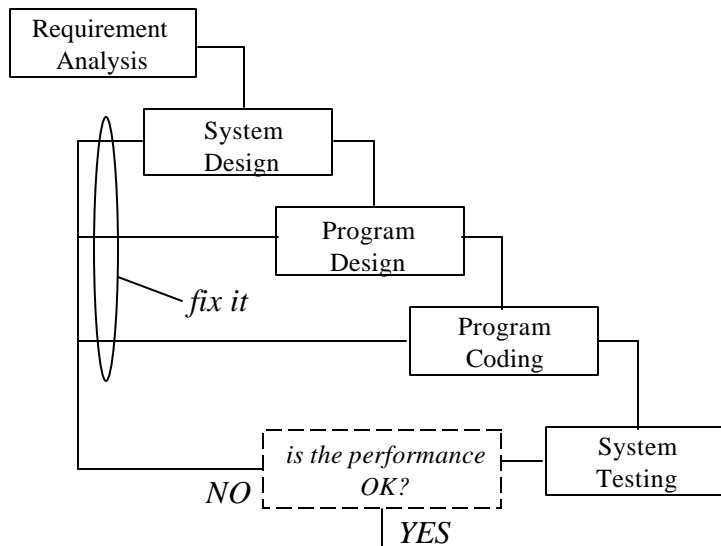
- *service levels (response times, throughputs, etc)*
- *hardware/software base (client and software platforms, networking technologies, DBMSs)*



27

© 2004 D. A. Menascé. All Rights Reserved.

Traditional Software Development Life Cycle



28

© 2004 D. A. Menascé. All Rights Reserved.

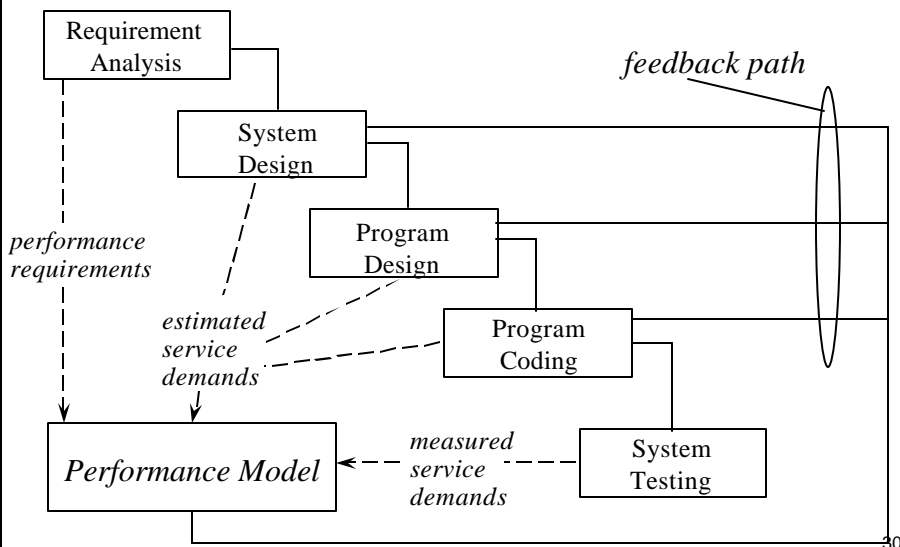
Traditional Software Development Life Cycle

- Common approach:
 - consider Functional Requirements only during development and check Performance Requirements at the end.
 - fix the system if performance is not good!
- Problem:
 - it is very costly and time consuming to fix the problem after the system is ready!
 - fixing the problem may imply in major software rewrites.

29

© 2004 D. A. Menascé. All Rights Reserved.

Integrating Software Performance Engineering Into the Software Development Life Cycle



© 2004 D. A. Menascé. All Rights Reserved.

Service Demand Generation

$$D_{i,r} = \sum_{s \in S_{i,r}} n_s \times p_s \times D_{i,r}^s$$

where,

$D_{i,r}^s$: average service demand at device i for class r due to statement s .

$S_{i,r}$: set of statements that generate demands at device i for class r .

n_s : average number of time that statement s is executed.

p_s : probability that statement s is executed.

31