On the Use of Performance Models to Design Self-Managing Computer Systems

Daniel A. Menascé
and
Mohamed Bennani
Dept. of Computer Science
George Mason University
www.cs.gmu.edu/faculty/menasce.html

A Computer System

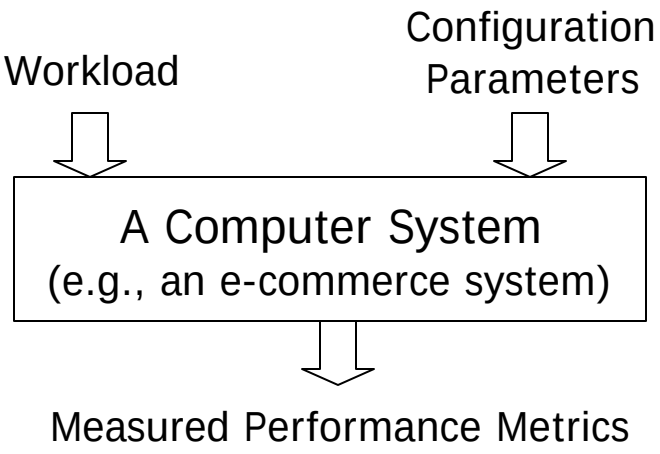

A Performance Model

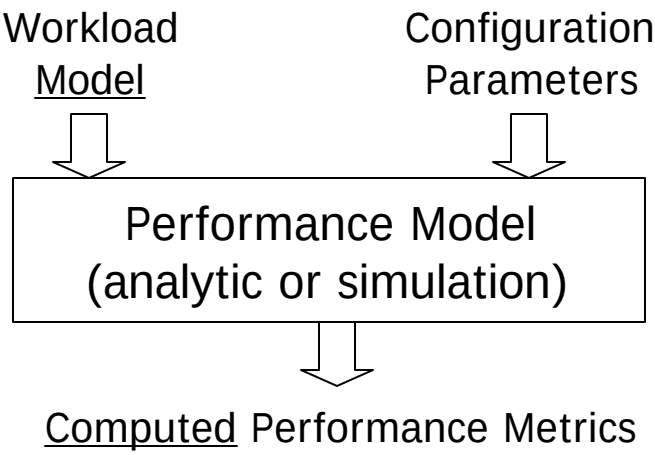

3

What are performance models good for?

- At the design stage:
  - Compare competing design alternatives.
    - A large number of low capacity servers vs. a small number of large capacity servers?

4

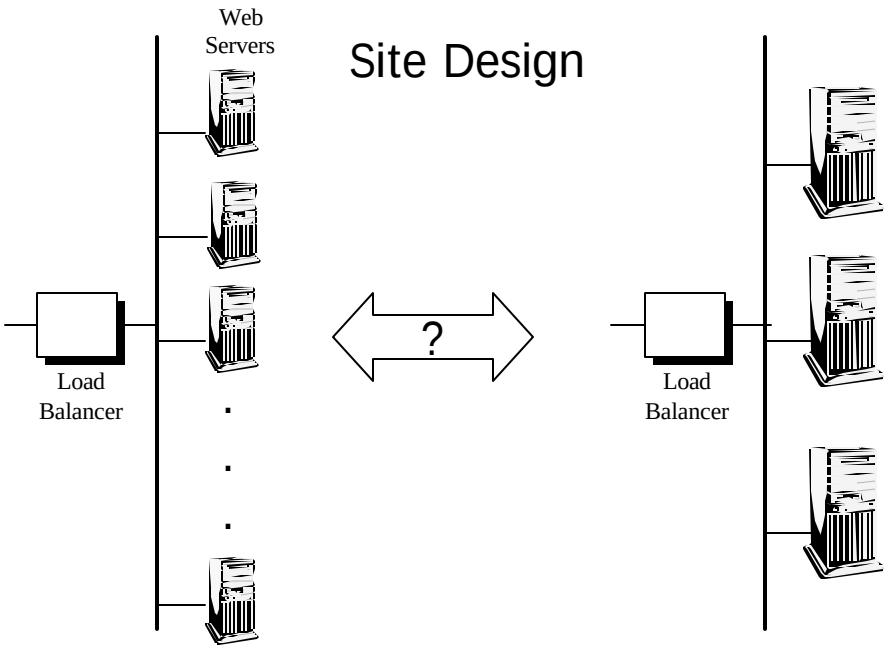

5

What are performance models good for?

- At the design stage:
  - Compare competing design alternatives.
    - A large number of low capacity servers vs. a small number of large capacity servers?
- During production:
  - Medium and long-term (weeks and months):
    - Capacity planning.

6

Capacity Planning

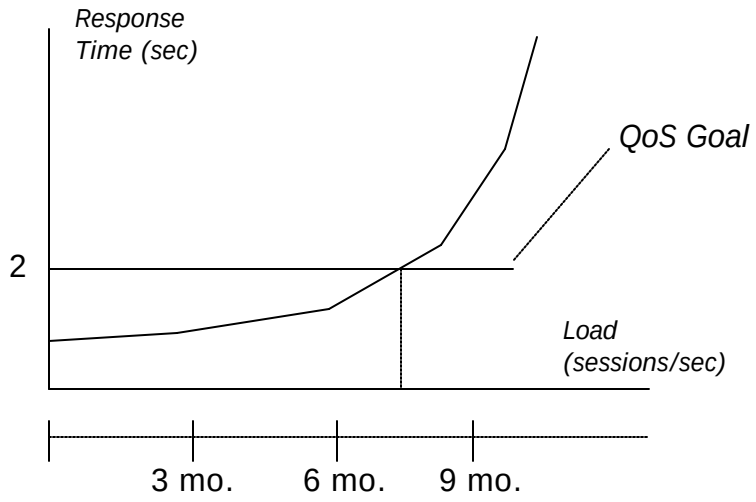

7

What are performance models good for?

- At the design stage:
  - Compare competing design alternatives.
    - A large number of low capacity servers vs. a small number of large capacity servers?
- During production:
  - Medium and long-term (weeks and months):
    - Capacity planning.
  - Short-term (minutes):
    - Dynamic reconfiguration.

8

Dynamic QoS Control: Motivation

- Computer systems are becoming very complex and composed of multiple tiers.

9

Multi-tier Architecture

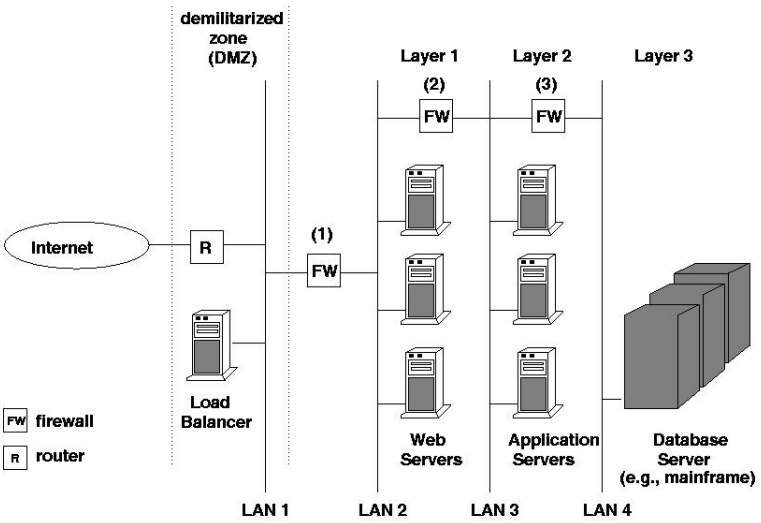

10

Dynamic QoS Control: Motivation

- Computer systems are becoming very complex and composed of multiple tiers.
- The workload presents short-term variations with high peak-to-average ratios.

Multi-scale workload variation

3600 sec

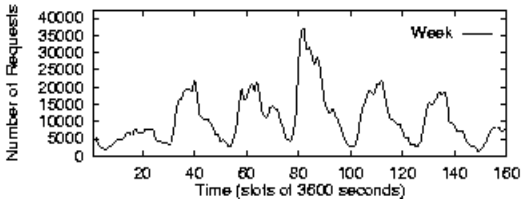

Multi-scale workload variation

3600 sec

60 sec

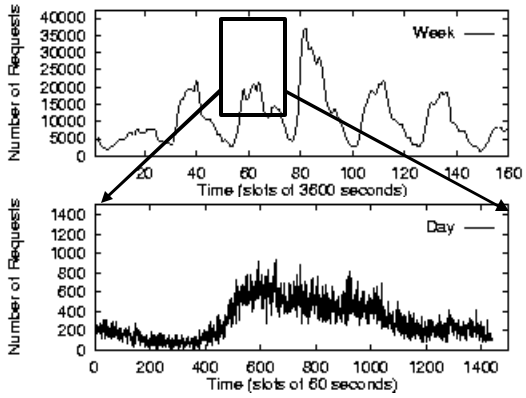

13

Multi-scale workload variation

3600 sec

60 sec

1 sec

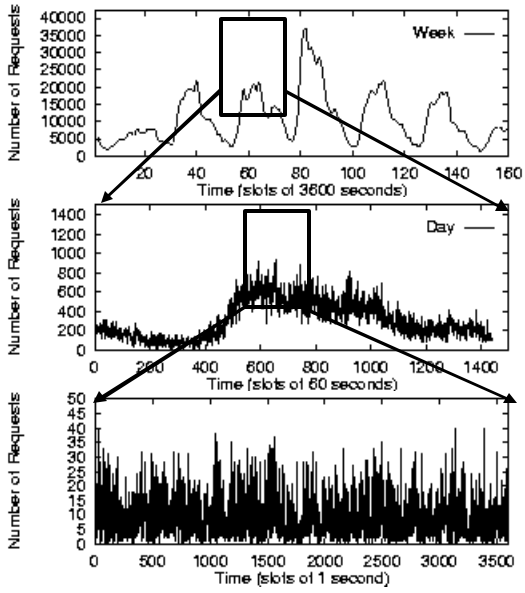

Dynamic QoS Control: Motivation

- E-commerce sites are complex and composed of multiple tiers.
- The workload presents short-term variations with high peak-to-average ratios.
- Many software and hardware parameters influence the performance of computer systems.

⇒ Manual reconfiguration is not an option.

Need self-managing systems!

15

Outline of the rest of the talk

- Basic Approach to Self-managing Systems
- QoS metrics
- A Multi-Threaded Server Example
- Simulation Experiments
- Experiments on an Actual Web Server
- Concluding Remarks

16

Outline of the rest of the talk

- Basic Approach to Self-managing Systems
- QoS metrics
- A Multi-Threaded Server Example
- Simulation Experiments
- Experiments on an Actual Web Server
- Concluding Remarks

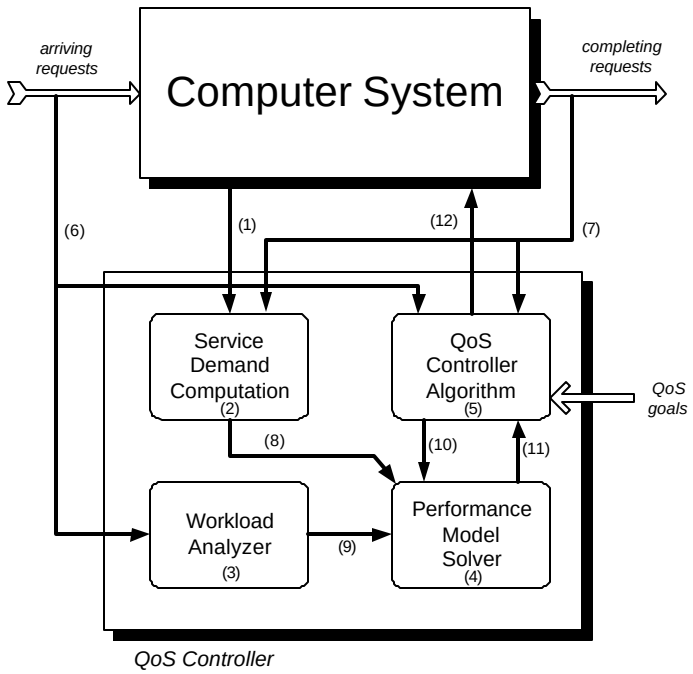

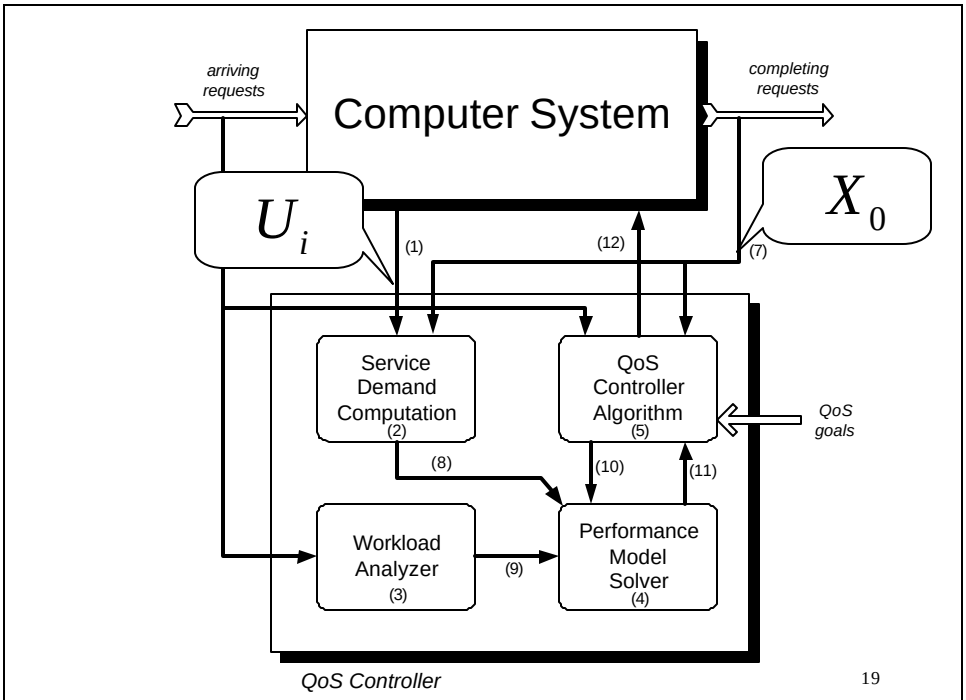

19

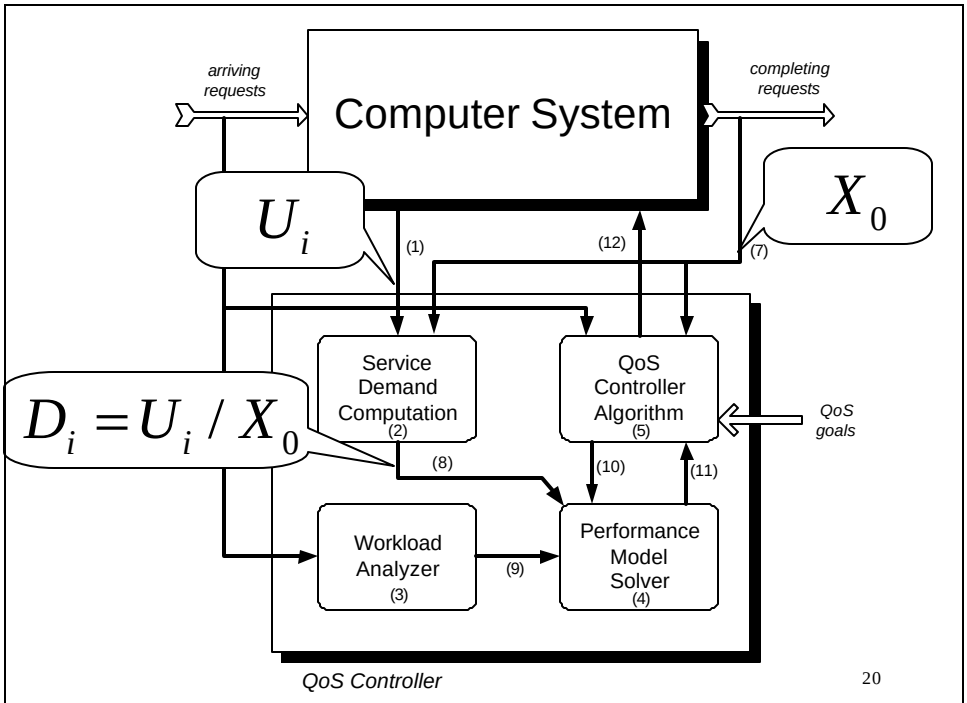

20

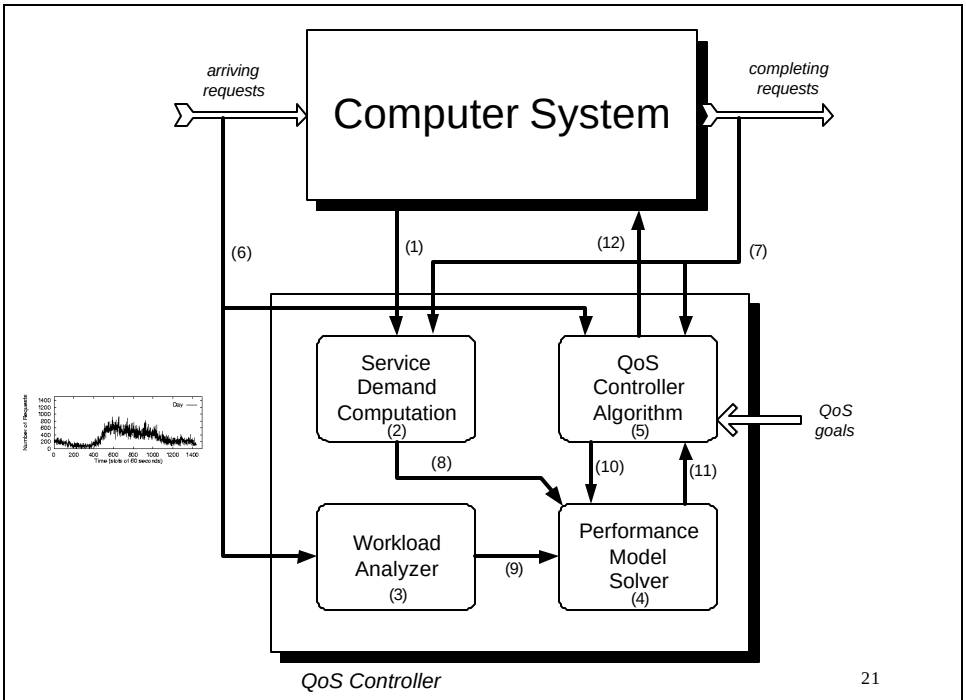

21

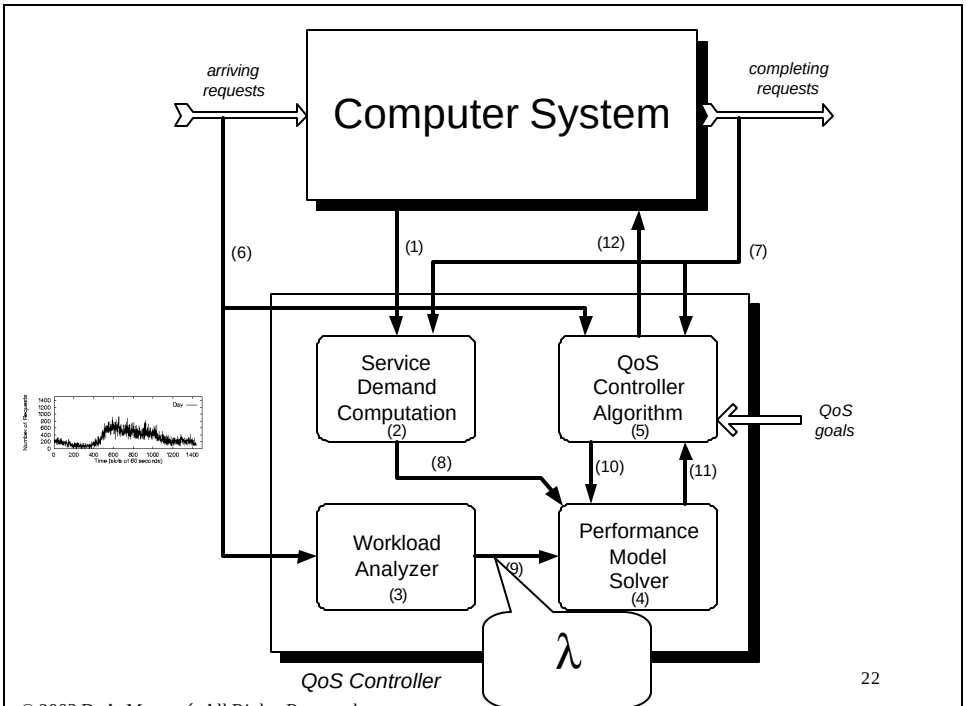

22

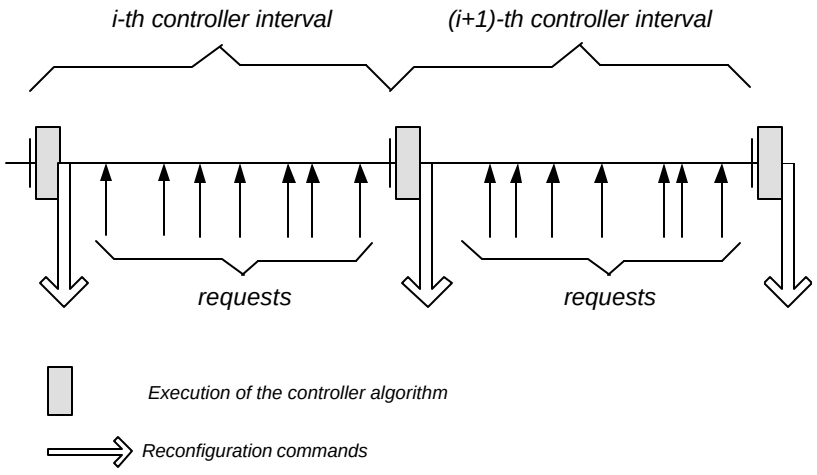

Heuristic Optimization Approach

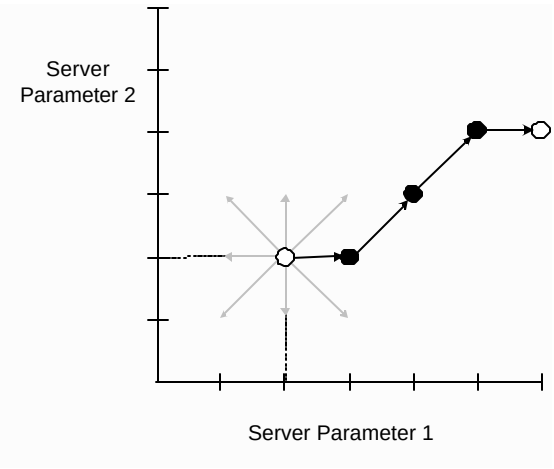

- The space of configuration points is searched using combinatorial search techniques.
- Each point has a QoS value computed through an analytic performance model.

Heuristic Optimization Approach

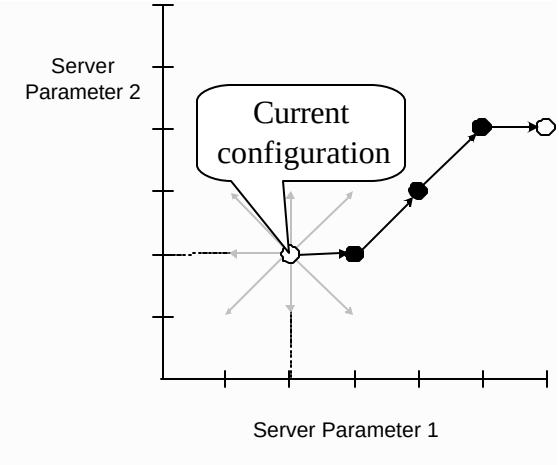

- The space of configuration points is searched using combinatorial search techniques.
- Each point has a QoS value computed through an analytic performance model.

25

Heuristic Optimization Approach

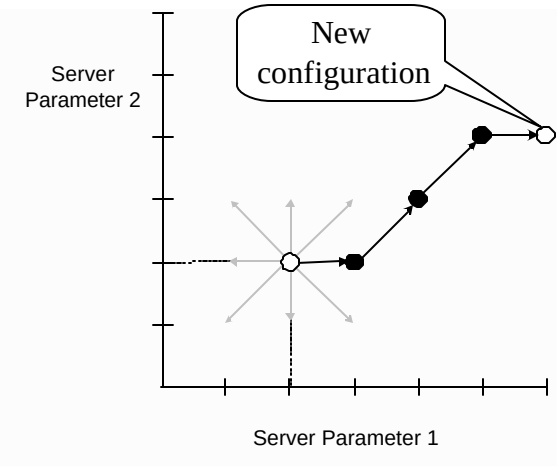

- The space of configuration points is searched using combinatorial search techniques.
- Each point has a QoS value computed through an analytic performance model.

26

Hill-Climbing Search

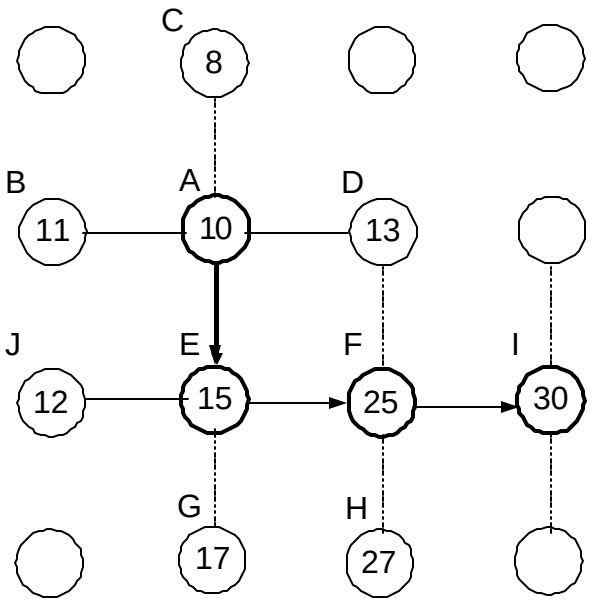

27

Hill-Climbing Search

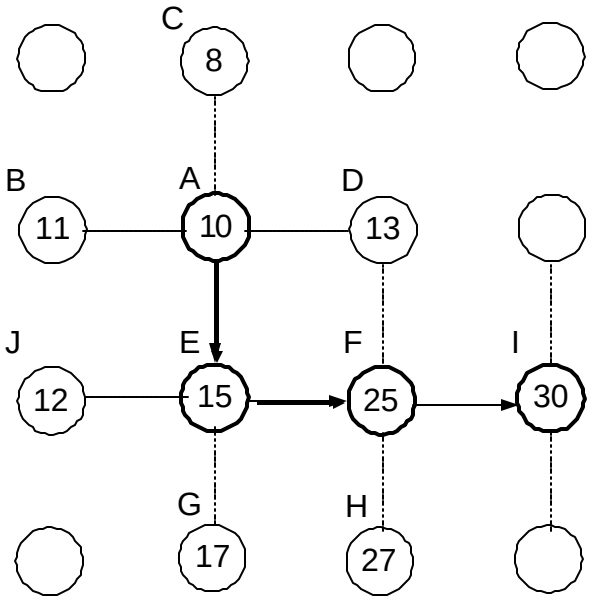

28

Hill-Climbing Search

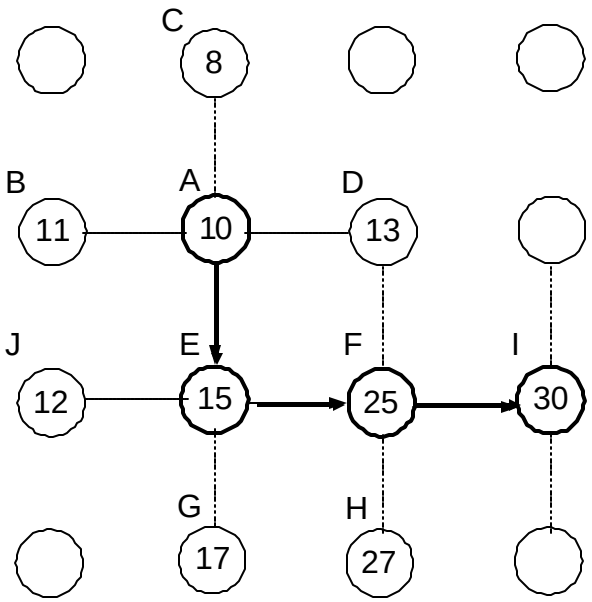

29

Problems with Hill-Climbing

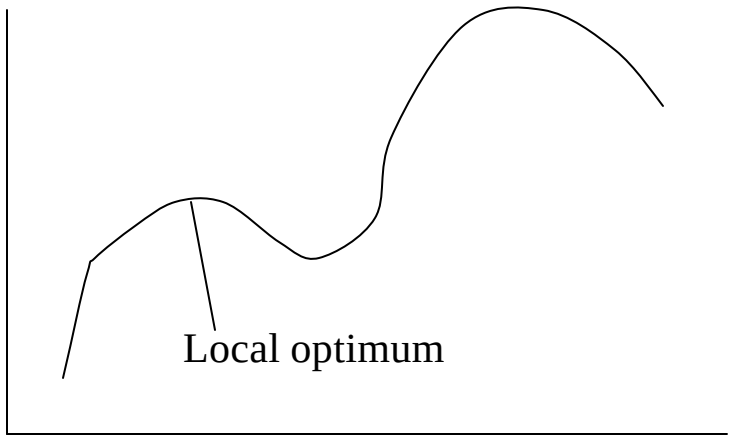

30

Beam Search

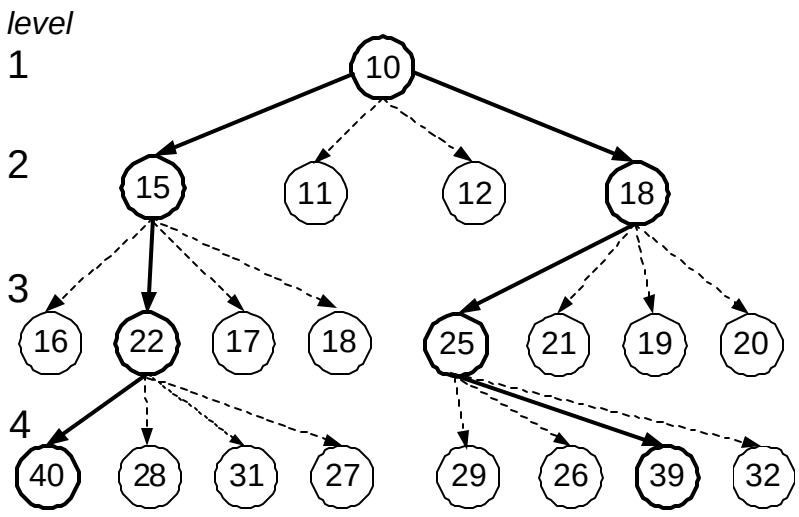

31

If I Had Used Hill Climbing Instead of Beam Search …

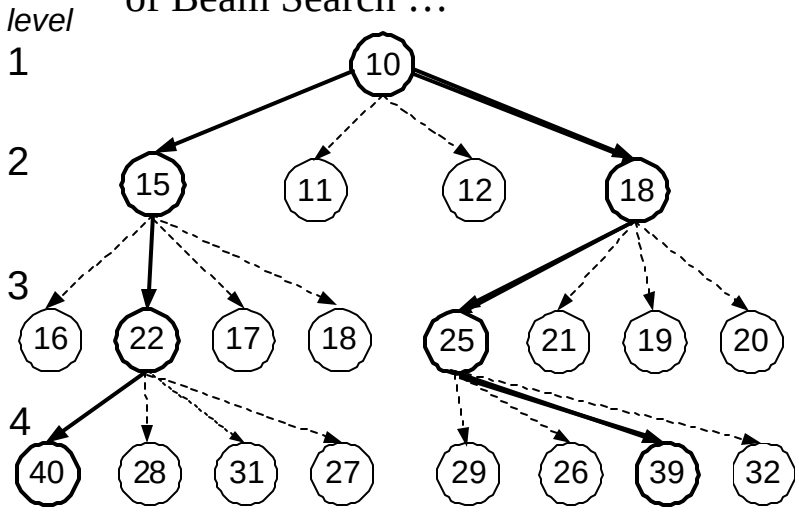

32

Outline of the rest of the talk

- Basic Approach to Self-managing Systems
- QoS metrics
- A Multi-Threaded Server Example
- Simulation Experiments
- Experiments on an Actual Web Server
- Concluding Remarks

QoS Metric

$$QoS = w_R \times \Delta QoS_R + w_P \times \Delta QoS_P + w_X \times \Delta QoS_X$$

w_R, w_P , and w_X are relative weights that indicate the relative importance of response time, throughput, and probability of rejection.

QoS Metric

$$QoS = w_R \times \Delta QoS_R + w_P \times \Delta QoS_P + w_X \times \Delta QoS_X$$

w_R, w_P , and w_X are relative weights that indicate the relative importance of response time, throughput, and probability of rejection.

ΔQoS_R, ΔQoS_P , and ΔQoS_X are relative deviations of the response time, throughput, and probability of rejection metrics with respect to their desired levels.

35

QoS Metric

$$QoS = w_R \times \Delta QoS_R + w_P \times \Delta QoS_P + w_X \times \Delta QoS_X$$

w_R, w_P , and w_X are relative weights that indicate the relative importance of response time, throughput, and probability of rejection.

ΔQoS_R, ΔQoS_P , and ΔQoS_X are relative deviations of the response time, throughput, and probability of rejection metrics with respect to their desired levels.

The QoS metric is a dimensionless number in the interval [-1, 1].
If all metrics meet or exceed their QoS targets than QoS = 0.

36

Response Time Deviation

$$\Delta QoS_R = \frac{R_{\max} - R_{measured}}{\max(R_{\max}, R_{measured})}$$

- = 0 if the response time meets its target.
- > 0 if the response time exceeds its target.
- < 0 if the response time does not meet its target.

- $\Delta QoS_R \le 1 - (\sum_{i=1}^{K} D_i) / R_{\max} < 1$
- $-1 < -(1 - R_{\max} / R_{measured}) \le \Delta QoS_R$

Probability of Rejection Deviation

$$\Delta QoS_P = \frac{P_{\max} - P_{measured}}{\max(P_{\max}, P_{measured})}$$

- = 0 if the probability of rejection meets its target.
- > 0 and = 1 if the probability of rejection exceeds its target.
- < 0 and = -1 if the probability of rejection does not meet its target.

Throughput Deviation

$$\Delta QoS_X = \frac{X_{measured} - X^*_{\min}}{\max(X_{measured}, X^*_{\min})}$$

- $X^*_{\min} = \min(l, X_{\min})$
- = 0 if the throughput meets its target.
- > 0 and < 1 if the throughput exceeds its target.
- < 0 and > -1 if the throughput does not meet its target.

39

Outline of the rest of the talk

- Basic Approach to Self-managing Systems
- QoS metrics
- A Multi-Threaded Server Example
- Simulation Experiments
- Experiments on an Actual Web Server
- Concluding Remarks

40

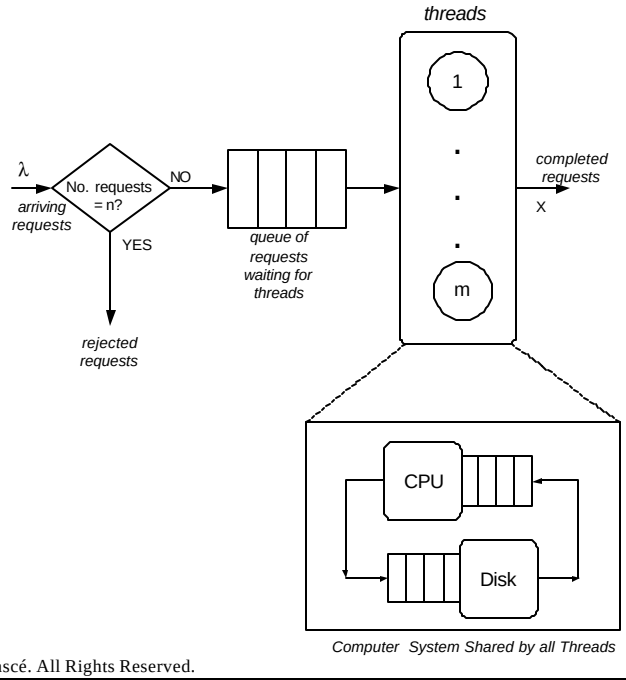

41

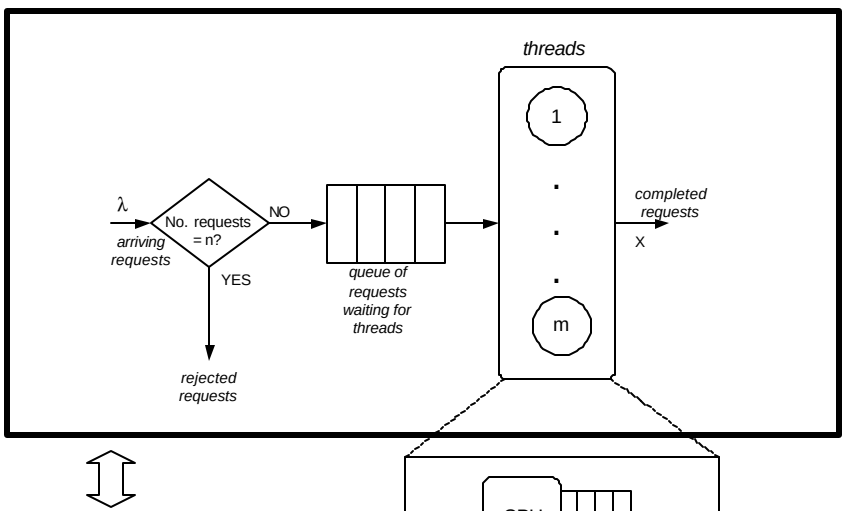

Markov Chain model used to find the probability distribution of the number of requests in the system.

42

Markov Chain Model

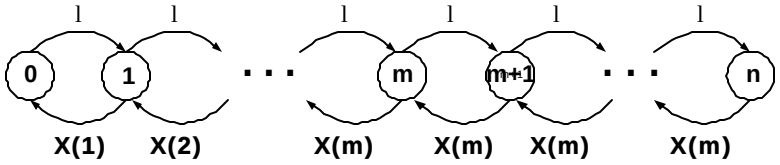

$$P_k = \begin{cases} P_0 \, l^{\,k} / b(k) & k = 1,\ldots,m \\ P_0 \, r^{\,k} \, X(m)^m / b(m) & k = m+1,\ldots,n \end{cases}$$

where $b(k) = X(1) \times X(2) \times \ldots \times X(k)$, $r = l / X(m)$, and

$$P_0 = \left[1 + \sum_{k=1}^{m} \frac{l^{\,k}}{b(k)} + \frac{r \times l^{\,m} (1 - r^{\,n-m})}{b(m)(1-r)} \right]^{-1}.$$

43

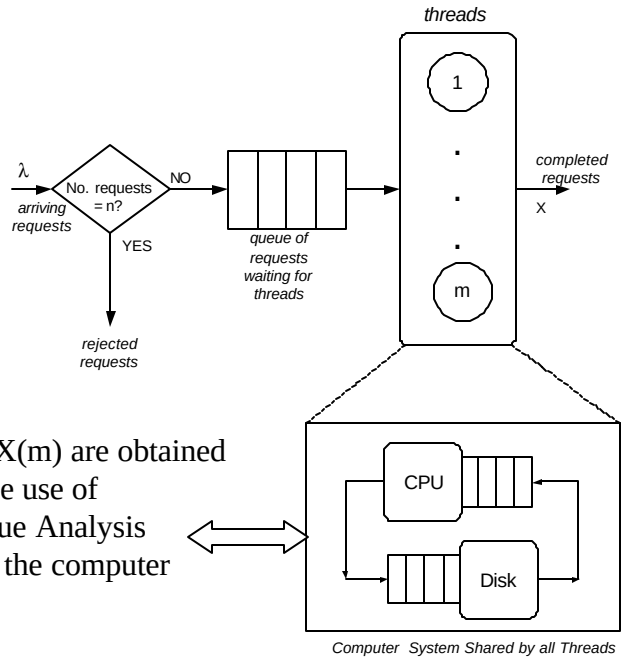

X(1), …, X(m) are obtained through the use of Mean Value Analysis applied to the computer system.

44

Outline of the rest of the talk

- Basic Approach to Self-managing Systems
- QoS metrics
- A Multi-Threaded Server Example
- Simulation Experiments
- Experiments on an Actual Web Server
- Concluding Remarks

Evolution of Workload Intensity

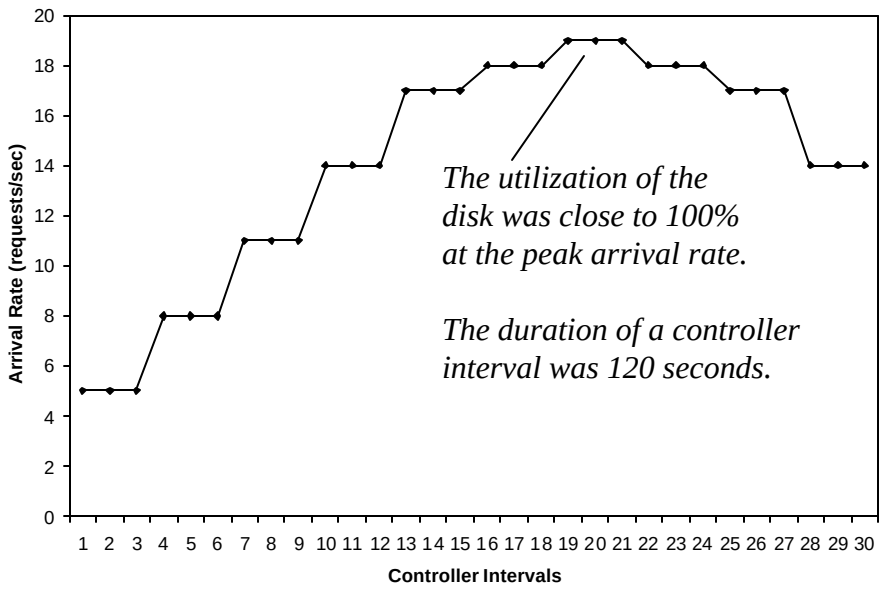

QoS Variation During the Experiment

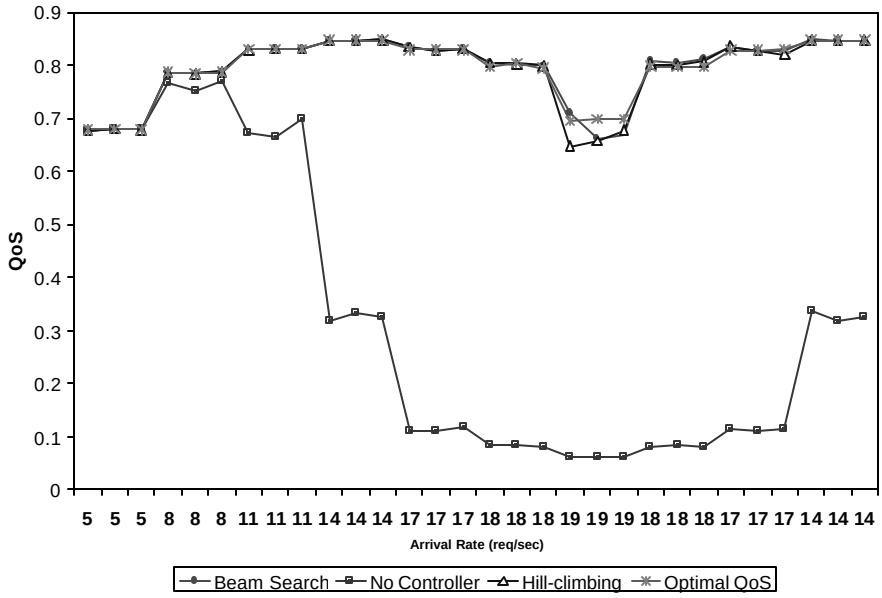

47

Some details about this experiment

- The "optimal" QoS was computed off-line by considering all possible 9,801 configurations at each controller interval.
- Both heuristics evaluated no more than 120 configurations (i.e, 1.2% of the total) per controller interval

48

Disk Utilization During the Experiment

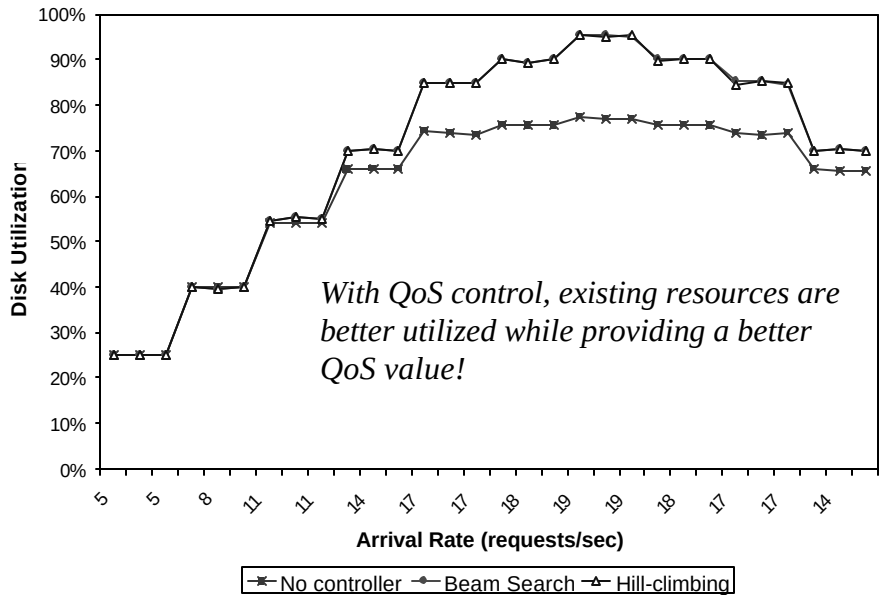

With QoS control, existing resources are better utilized while providing a better QoS value!

49

Response Time Variation During the Experiment

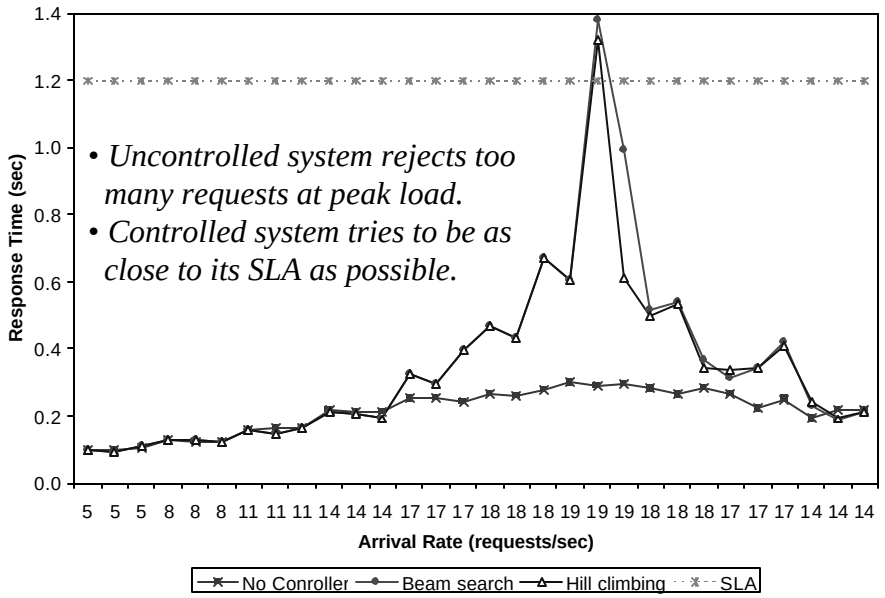

- *Uncontrolled system rejects too many requests at peak load.*
- *Controlled system tries to be as close to its SLA as possible.*

50

Throughput Variation During the Experiment

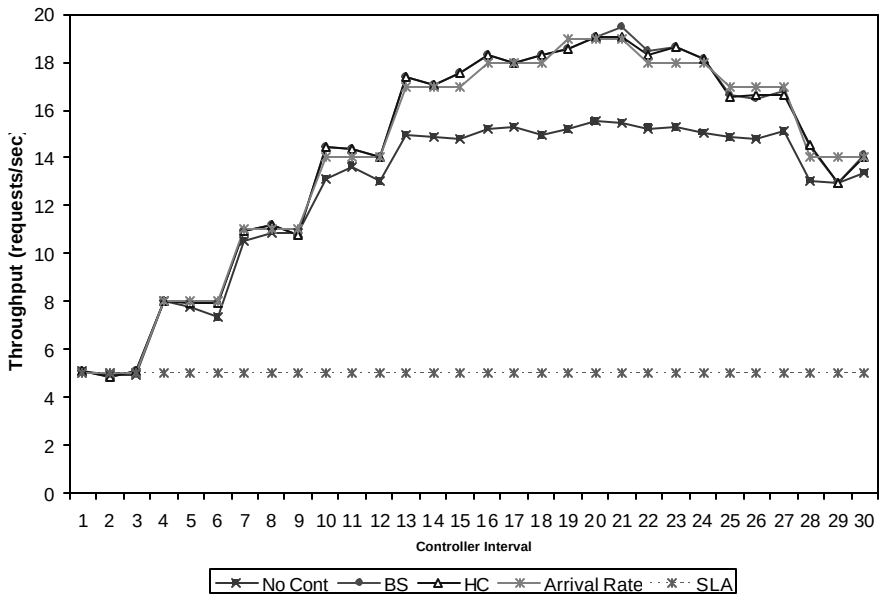

51

Outline of the rest of the talk

- Basic Approach to Self-managing Systems
- QoS metrics
- A Multi-Threaded Server Example
- Simulation Experiments
- Experiments on an Actual Web Server
- Concluding Remarks

52

Experiments with an Actual Web Server

- Controller code is the same as in the simulation.
- HTTP server: Apache 1.3.12 modified to allow for dynamic number of threads (m) and maximum number of requests (n).
- Workload generated by SURGE (Scalable URL Reference Generator)

53

Experiments with an Actual Web Server

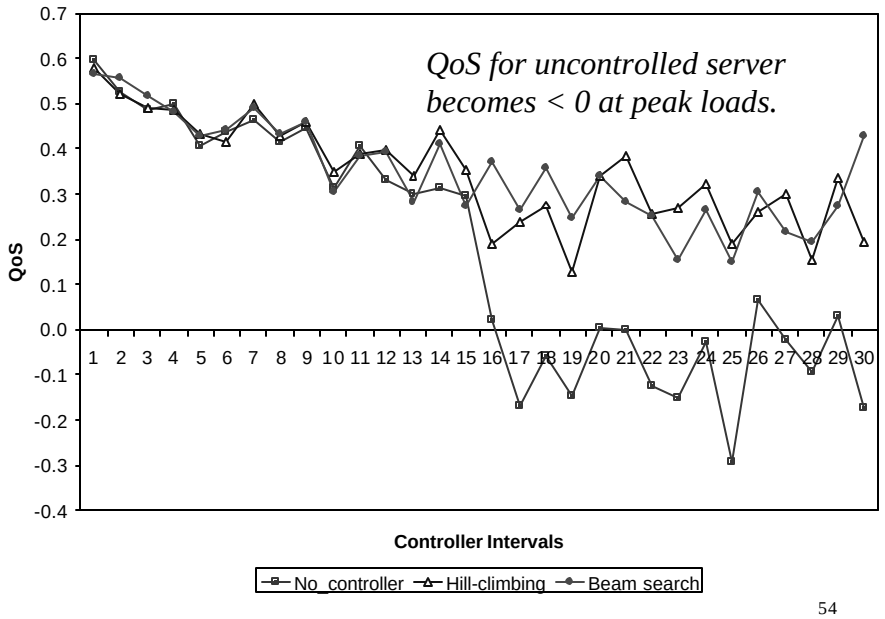

54

Outline of the rest of the talk

- Basic Approach to Self-managing Systems
- QoS metrics
- A Multi-Threaded Server Example
- Simulation Experiments
- Experiments on an Actual Web Server
- Concluding Remarks

Concluding Remarks

- Analytic models can be used to design self-managing systems due to their ability to track the evolution of the performance metrics as the workload and configuration parameters change.
- Experiments show that this approach is robust even when the coefficients of variation of the inter-arrival times and service times are high (around 3).

Highly Variable Workloads
Ca = 2 and Cs = 4

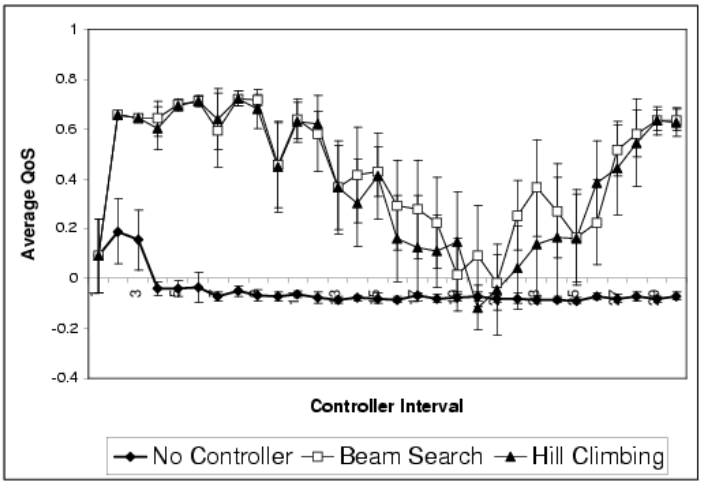

57

Concluding Remarks

- Self-managing systems provide better resource utilization while meeting the SLAs.
- Self-managing systems allow existing resources to be used to their fullest extent before upgrades are needed. Better ROI!
- Self-managing systems reduce personnel costs.
- Other approaches: control theory.

58

Bibliography

- "Assessing the Robustness of Self-Managing Computer Systems under Highly Variable Workloads," M. Bennani and D.A. Menascé, *Proc. International Conf. Autonomic Computing (ICAC-04)*, New York, NY, May 17-18, 2004.
- "Automatic QoS Control," D.A. Menascé, *IEEE Internet Computing*, January/February 2003, Vol. 7, No. 1.
- Preserving QoS of E-commerce Sites Through Self-Tuning: A Performance Model Approach," D. A. Menascé, R. Dodge and D. Barbara, *Proc. 2001 ACM Conference on E-commerce*, Tampa, FL, October 14-17, 2001.