

# Using Performance Models to Design Self-Configuring and Self-Optimizing Computer Systems

---

Prof. Daniel Menascé  
Department of Computer Science  
E-Center for E-Business  
George Mason University  
Fairfax, VA, USA  
[Menasce@cs.gmu.edu](mailto:Menasce@cs.gmu.edu)  
[www.cs.gmu.edu/faculty/menasce.html](http://www.cs.gmu.edu/faculty/menasce.html)

© 2004 D. A. Menascé. All Rights Reserved.

1



2



Huge number of devices  
 Huge number of data sources  
 Many different data formats  
 Heterogeneous devices  
 Widely varying capacity  
 Wired and wireless  
 Widely varying QoS requirements

3



Node failures  
 Connectivity failures  
 Security attacks  
 Limited battery power

4

## Characteristics of the new generation of distributed software systems

- ☐ Highly distributed
- ☐ Component-based (for reusability)
- ☐ Service-oriented architectures (SOA)
- ☐ Unattended operation
- ☐ Hostile environments
- ☐ Composed of a large number of “replaceable” components discovered at run-time
- ☐ Run on a multitude of (unknown and heterogeneous) hardware and network platforms

## Requirements of Next Generation of Large Distributed Systems

- ☐ Adaptable and self-configurable to changes in workload intensity:
  - QoS requirements at the application and component level must be met.
- ☐ Adaptable and self-configurable to withstand attacks and failures:
  - Availability and security requirements must be met.



self-configurable, self-optimizing, self-healing,  
and self-protecting

# Important Technologies

- ❑ Web Services:
  - SOAP, UDDI, WSDL
- ❑ Grid Computing
- ❑ Peer to Peer Networks
- ❑ Wireless Networking
- ❑ Sensor and ad-hoc networks

# Challenges

- ❑ Dynamically changing application structure.
- ❑ Hard to characterize the workload.
  - unpredictable
  - dynamically changing services
  - application adaptation
- ❑ Difficult to build performance models.
  - moving target
- ❑ Multitude of QoS metrics at various levels of a distributed architecture
  - response time, jitter, throughput, availability, survivability, recovery time after attack/failure, call drop rate, access failure rate, packet delay, packet drop rate.
- ❑ Tradeoffs between QoS metrics (response time vs. availability, response time vs. security)

## Challenges (cont'd)

- ❑ Need to perform transient, critical time (e.g., terrorism attack or catastrophic failures) analysis of QoS compliance. Steady-state analysis is not enough.
- ❑ Mapping of global SLAs to local SLAs
  - Cost and pricing issues.
- ❑ QoS monitoring, negotiation, and enforcement.
- ❑ Platform-neutral representation of QoS goals and contracts.
- ❑ Resource management: resource reservation, resource allocation, admission control.
  - non-dedicated resources

## What we need ...

- ❑ Design self-regulating (autonomic systems)
- ❑ Embed within each system component:
  - Monitoring capabilities
  - Negotiation capabilities (requires predictive modeling power)
  - Self-protection and recovery capabilities (for attacks, failures, and overloads)
- ❑ Push more of the desired system features to individual components
- ❑ Design QoS-aware middleware
  - QoS negotiation protocols
  - Mapping of global to local SLAs
  - QoS monitoring

## Rest of this talk ...

- ❑ Novel uses for performance models
- ❑ Two examples of self-regulating systems:
  - A three-tiered e-commerce system
  - QoS-aware software components
- ❑ Concluding Remarks

## Rest of this talk ...

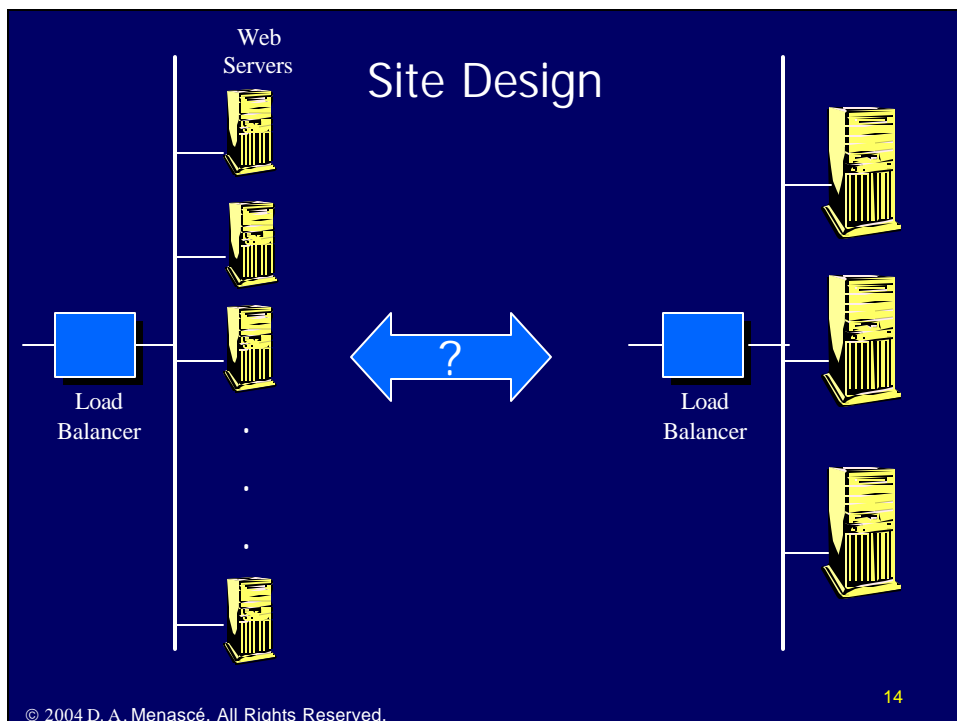
- ❑ Novel uses for performance models
- ❑ Two examples of self-regulating systems:
  - A three-tiered e-commerce system
  - QoS-aware software components
- ❑ Concluding Remarks

# What are performance models good for?

## □ At the design stage:

### ➤ Compare competing design alternatives.

- A large number of low capacity servers vs. a small number of large capacity servers?



# What are performance models good for?

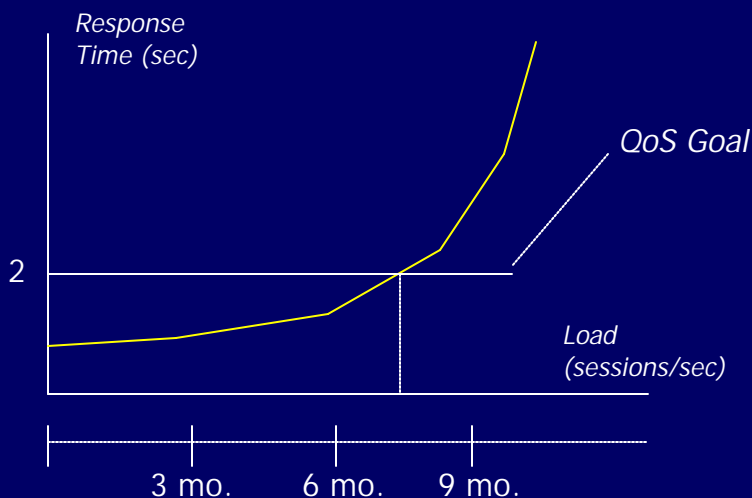
## □ At the design stage:

- Compare competing design alternatives.
  - A large number of low capacity servers vs. a small number of large capacity servers?

## □ During production:

- Medium and long-term (weeks and months):
  - Capacity planning.

# Capacity Planning





# What are performance models good for?

- ❑ At the design stage:
  - Compare competing design alternatives.
    - A large number of low capacity servers vs. a small number of large capacity servers?
- ❑ During production:
  - Medium and long-term (weeks and months):
    - Capacity planning.
  - Short-term (minutes):
    - Dynamic reconfiguration.

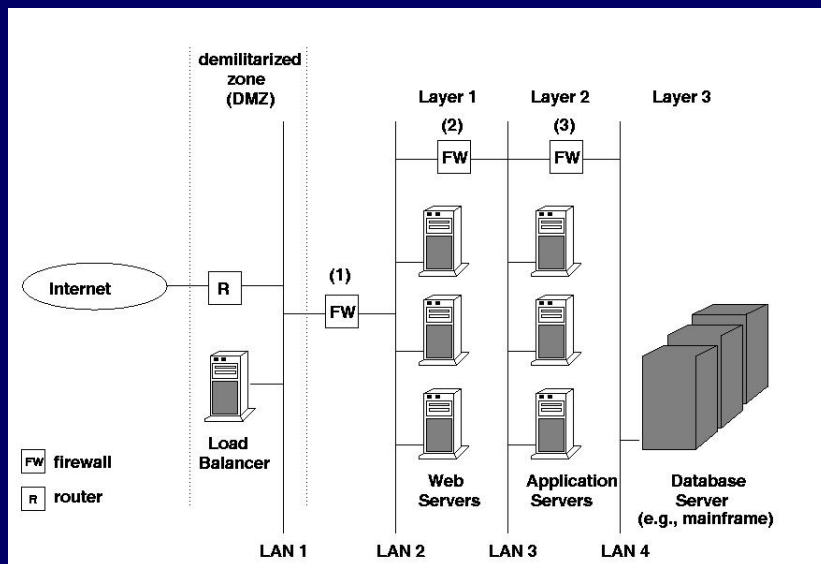
## Rest of this talk ...

- ❑ Novel uses for performance models
- ❑ Two examples of self-regulating systems:
  - A three-tiered e-commerce system
  - QoS-aware software components

# Automatic QoS Control: Motivation

- ❑ Modern computer systems are complex and composed of multiple tiers.

## Multi-tier Architecture



# Automatic QoS Control: Motivation


- ❑ Modern computer systems are complex and composed of multiple tiers.
- ❑ The workload presents short-term variations with high peak-to-average ratios.

## Multi-scale time workload variation

3600 sec


60 sec

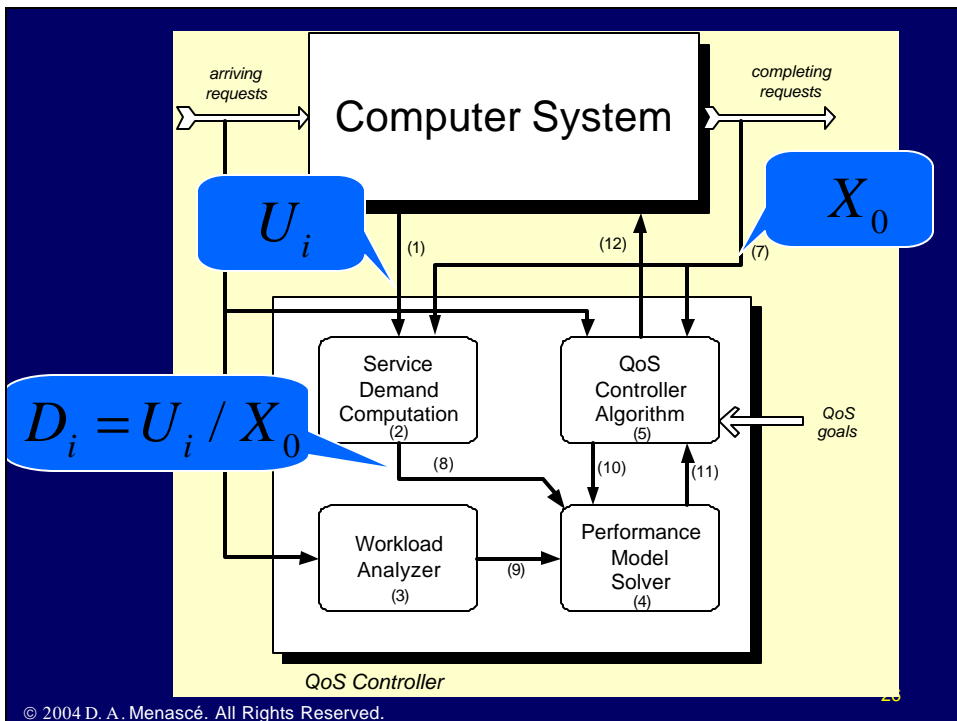
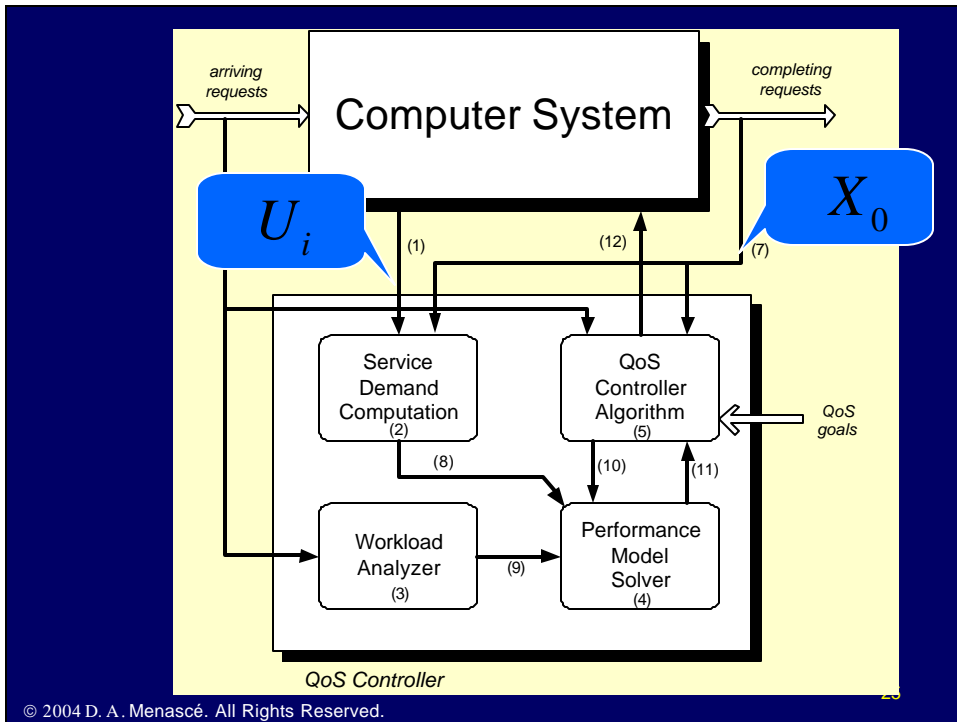
1 sec

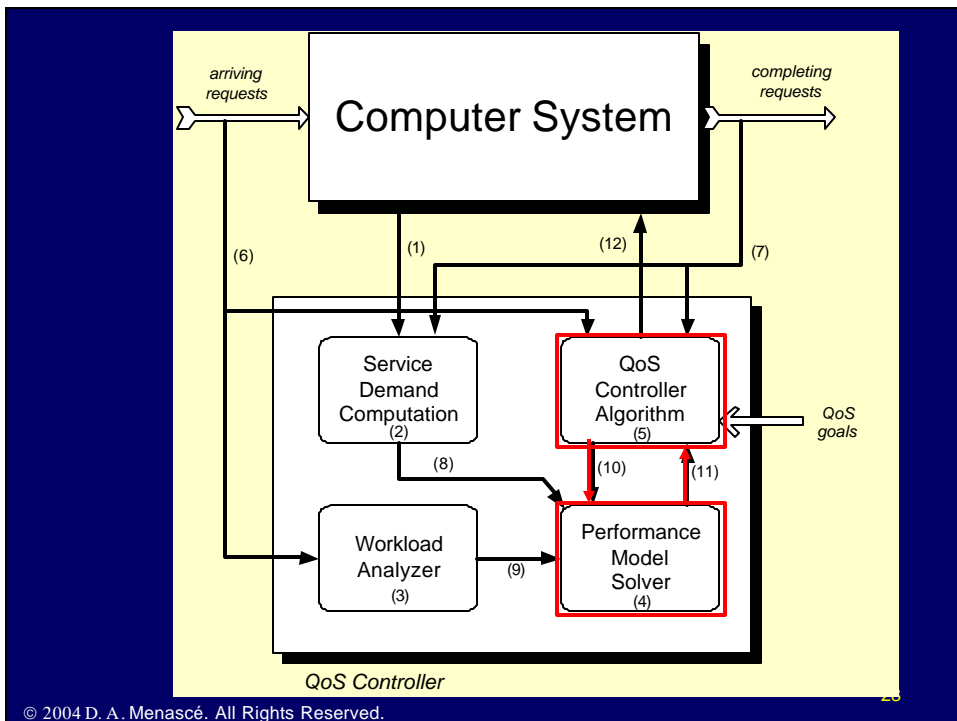
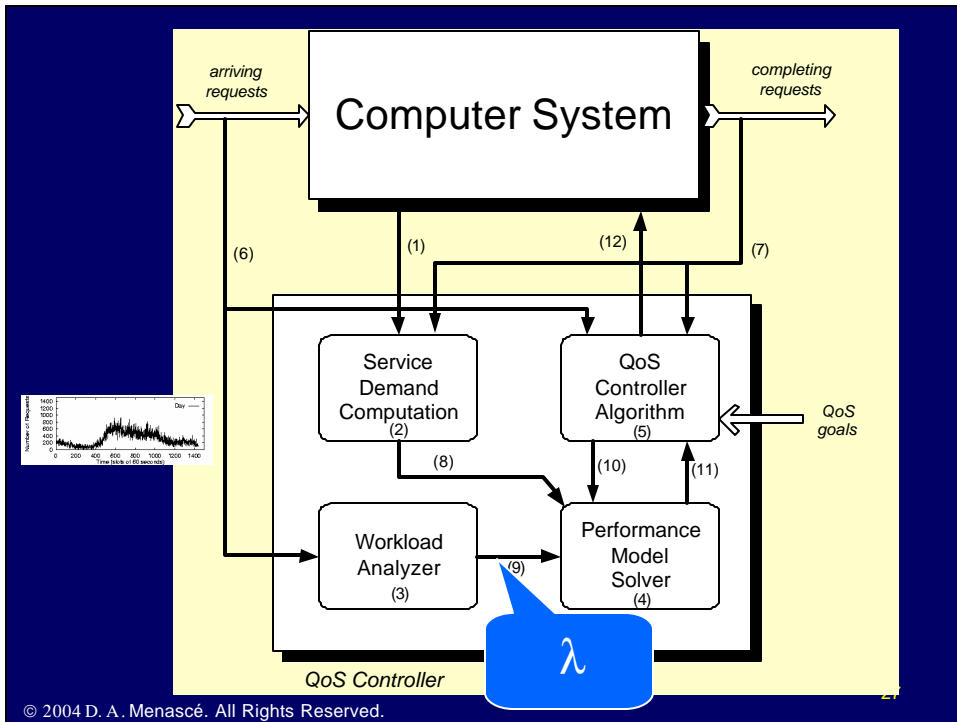


# Automatic QoS Control: Motivation


- ❑ Modern computer systems are complex and composed of multiple tiers.
  - ❑ The workload presents short-term variations with high peak-to-average ratios.
  - ❑ Many software and hardware parameters influence the performance of e-commerce sites.
- ➔ Manual reconfiguration is not an option!  
Need self-managing systems.







## Controller Interval



© 2004 D. A. Menascé. All Rights Reserved.

29

## Combined QoS Metric

$$QoS = w_R \times \Delta QoS_R + w_P \times \Delta QoS_P + w_X \times \Delta QoS_X$$

$w_R$ ,  $w_P$ , and  $w_X$  are relative weights that indicate the relative importance of response time, throughput, and probability of rejection.

© 2004 D. A. Menascé. All Rights Reserved.

30

## QoS Metric

$$QoS = w_R \times \Delta QoS_R + w_P \times \Delta QoS_P + w_X \times \Delta QoS_X$$

$w_R$ ,  $w_P$ , and  $w_X$  are relative weights that indicate the relative importance of response time, throughput, and probability of rejection.

⇒  $\Delta QoS_R$ ,  $\Delta QoS_P$ , and  $\Delta QoS_X$  are relative deviations of the response time, throughput, and probability of rejection metrics with respect to their desired levels.

## QoS Metric

$$QoS = w_R \times \Delta QoS_R + w_P \times \Delta QoS_P + w_X \times \Delta QoS_X$$

$w_R$ ,  $w_P$ , and  $w_X$  are relative weights that indicate the relative importance of response time, throughput, and probability of rejection.

$\Delta QoS_R$ ,  $\Delta QoS_P$ , and  $\Delta QoS_X$  are relative deviations of the response time, throughput, and probability of rejection metrics with respect to their desired levels.

⇒ The QoS metric is a dimensionless number in the interval [-1, 1].



## QoS Metric

$$QoS = w_R \times \Delta QoS_R + w_P \times \Delta QoS_P + w_X \times \Delta QoS_X$$

$w_R$ ,  $w_P$ , and  $w_X$  are relative weights that indicate the relative importance of response time, throughput, and probability of rejection.

$\Delta QoS_R$ ,  $\Delta QoS_P$ , and  $\Delta QoS_X$  are relative deviations of the response time, throughput, and probability of rejection metrics with respect to their desired levels.

The QoS metric is a dimensionless number in the interval [-1, 1].

⇒ If all metrics meet or exceed their QoS targets, QoS = 0.

## Response Time Deviation

$$\Delta QoS_R = \frac{R_{\max} - R_{\text{measured}}}{\max(R_{\max}, R_{\text{measured}})}$$

- = 0 if the response time meets its target.
- > 0 if the response time exceeds its target.
- < 0 if the response time does not meet its target.

$$\Delta QoS_R \leq 1 - (\sum_{i=1}^K D_i) / R_{\max} < 1$$

$$-1 < -(1 - R_{\max} / R_{\text{measured}}) \leq \Delta QoS_R$$

## Probability of Rejection Deviation

$$\Delta QoS_P = \frac{P_{\max} - P_{\text{measured}}}{\max(P_{\max}, P_{\text{measured}})}$$


- = 0 if the probability of rejection meets its target.
- > 0 and = 1 if the probability of rejection exceeds its target.
- < 0 and = -1 if the probability of rejection does not meet its target.

## Throughput Deviation

$$\Delta QoS_X = \frac{X_{\text{measured}} - X_{\min}^*}{\max(X_{\text{measured}}, X_{\min}^*)}$$

- $X_{\min}^* = \min(I, X_{\min})$
- = 0 if the throughput meets its target.
- > 0 and < 1 if the throughput exceeds its target.
- < 0 and > -1 if the throughput does not meet its target.

# Heuristic Optimization Approach




- The space of configuration points is searched using a combinatorial search technique.
- Each point has a QoS value computed through an analytic performance model.

$$QoS = f(\vec{W}, c_1, c_2, \dots, c_m)$$

© 2004 D. A. Menascé. All Rights Reserved.

37

# Heuristic Optimization Approach




- The space of configuration points is searched using combinatorial search techniques.
- Each point has a QoS value computed through an analytic performance model.

© 2004 D. A. Menascé. All Rights Reserved.

38

# Heuristic Optimization Approach




- The space of configuration points is searched using combinatorial search techniques.
- Each point has a QoS value computed through an analytic performance model.

© 2004 D. A. Menascé. All Rights Reserved.

39


## Hill-Climbing Search



© 2003 D. A. Menascé. All Rights Reserved.

40


## Hill-Climbing Search



© 2003 D. A. Menascé. All Rights Reserved.

41


## Hill-Climbing Search



© 2003 D. A. Menascé. All Rights Reserved.

42

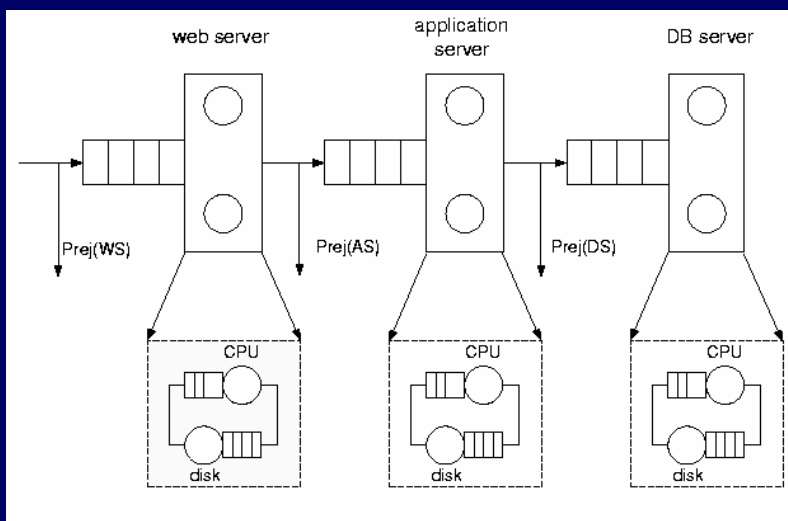
# Beam Search



© 2003 D. A. Menascé. All Rights Reserved.

43

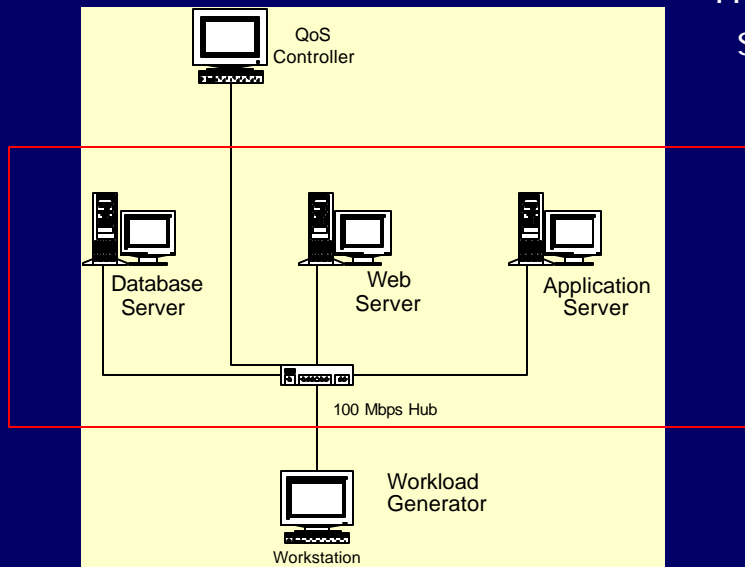
## A Queuing Model is Used to Compute QoS Values



© 2004 D. A. Menascé. All Rights Reserved.

44

# Prototype Configuration




TPC-W  
site

© 2004 D. A. Menascé. All Rights Reserved.

45

# Experiment Results

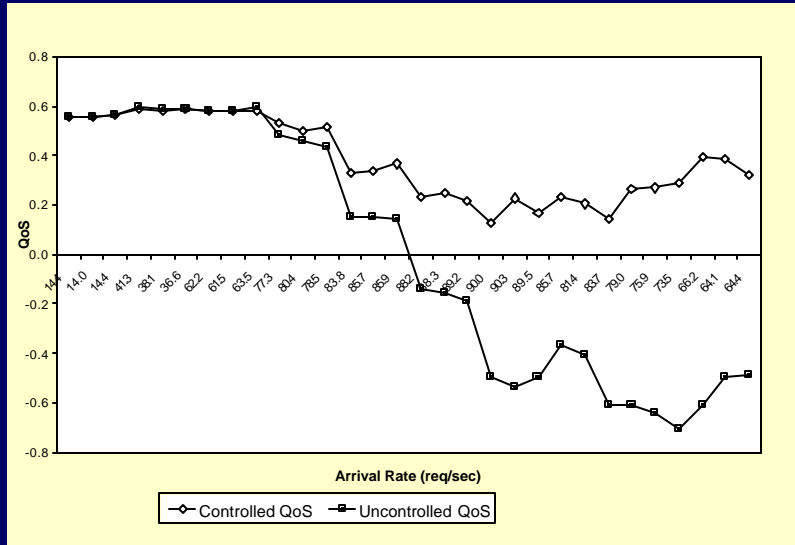
## Arrival rate



© 2004 D. A. Menascé. All Rights Reserved.

46

# Results of QoS Controller

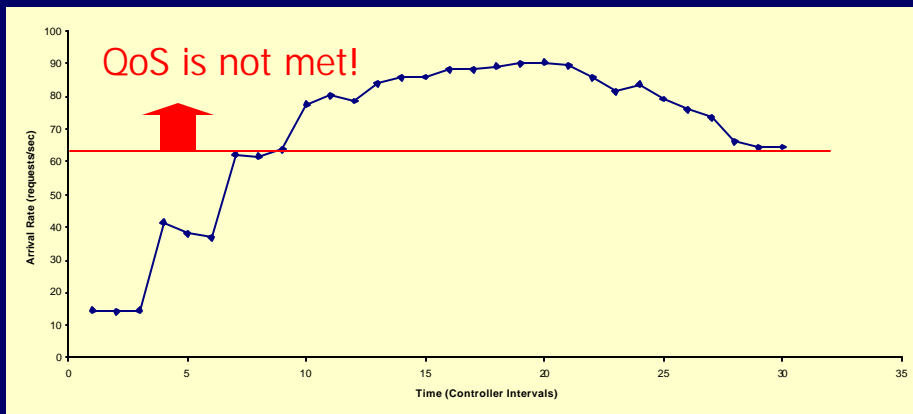


© 2004 D. A. Menascé. All Rights Reserved.

47

# Experiment Results

Arrival rate



© 2004 D. A. Menascé. All Rights Reserved.

48




## Variable inter-arrival and service times of requests

- ❑ Real workloads exhibit high variability in:
  - Traffic intensity
  - Service demands at various system resources
- ❑ Need to investigate the efficiency of the proposed self-managing technique under these conditions
- ❑ Consider variability in requests inter-arrival time and requests service times at physical resources (e.g., CPU, disk)


© 2004 D. A. Menascé. All Rights Reserved.

49


## Effect of Varying the COV of the Service Time ( $C_a = 1$ )



$C_s = 1$



$C_s = 2$





$C_s = 4$


© 2004 D. A. Menascé. All Rights Reserved.

50

## Effect of Varying the COV of the Interarrival Time ( $C_s = 1$ )



Ca = 1




Ca = 2

Ca = 4

51

© 2004 D. A. Menascé. All Rights Reserved.

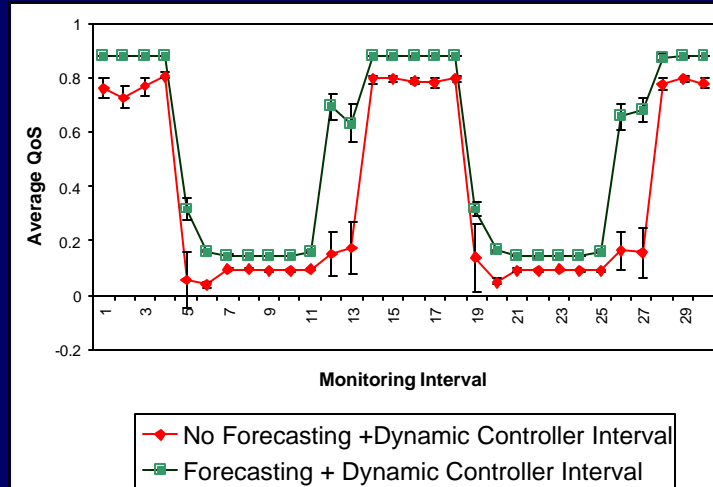
## Extreme values for Ca and Cs



© 2004 D. A. Menascé. All Rights Reserved.c

52

# Dynamic Controller Interval and Workload Forecasting



© 2004 D. A. Menascé. All Rights Reserved.c

53


## Sensitivity of Controller to SLAs

- ❑ Need to investigate the controller behavior in the case of a variation in the SLAs
- ❑ We ran experiments for stricter and more relaxed SLAs
  - Base:  $R_{max} = 1.2$  sec,  $X_{min} = 5$  req/sec,  $P_{max} = 0.05$
  - Strict:  $R_{max} = 1.0$  sec,  $X_{min} = 7$  req/sec,  $P_{max} = 0.03$
  - Relaxed:  $R_{max} = 1.5$  sec,  $X_{min} = 4$  req/sec,  $P_{max} = 0.10$
- ❑ Used  $C_a = C_s = 2$

© 2004 D. A. Menascé. All Rights Reserved.

54

## Sensitivity of Controller to SLAs



© 2004 D. A. Menascé. All Rights Reserved.

55

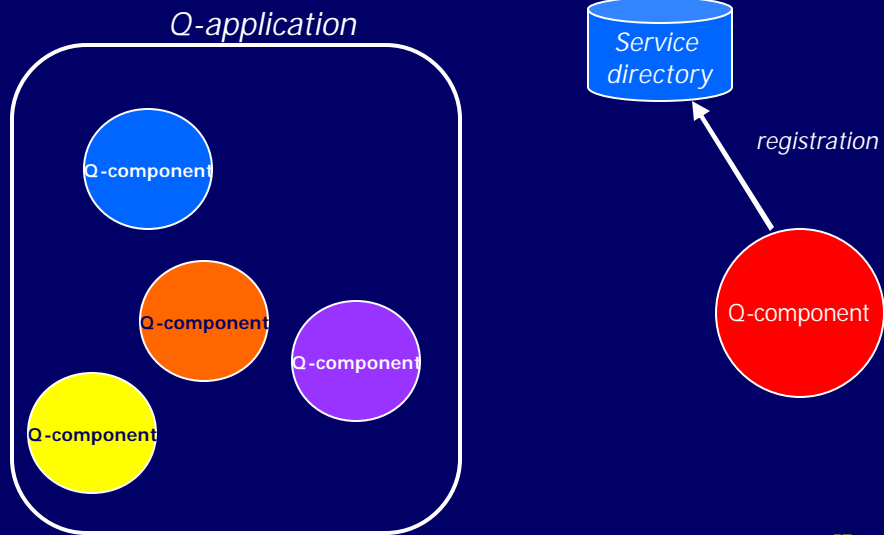
## Rest of this talk ...

- ❑ Novel uses for performance models
- ❑ Two examples of self-regulating systems:
  - A three-tiered e-commerce system
  - QoS-aware software components
- ❑ Concluding Remarks

© 2004 D. A. Menascé. All Rights Reserved.

56

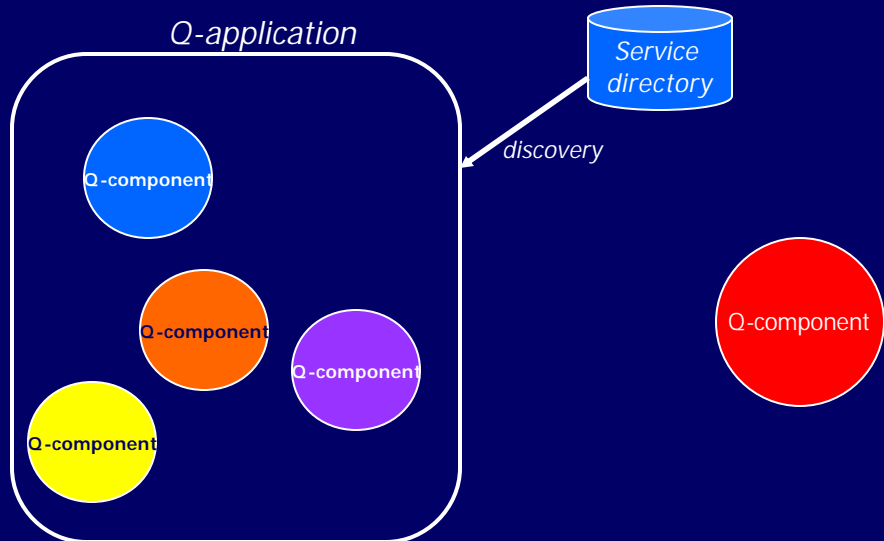
# Q-Applications and Q-components



© 2004 D. A. Menascé. All Rights Reserved.

57

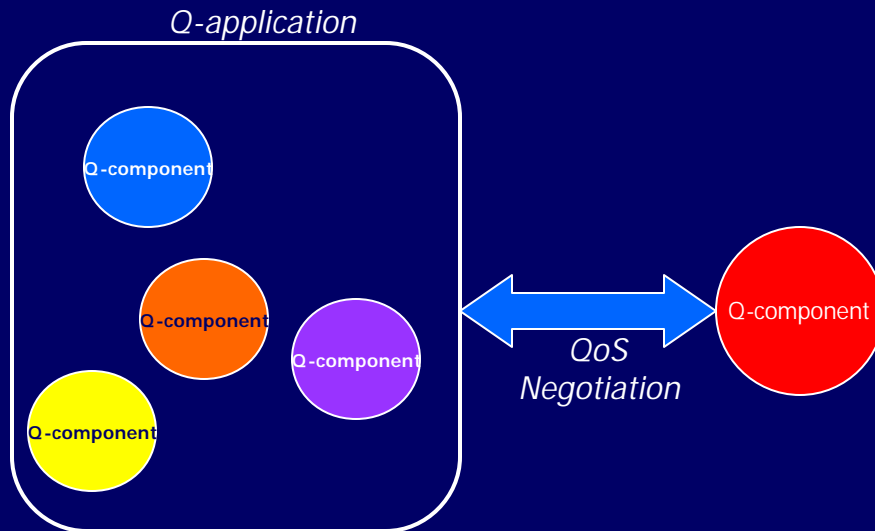
# Q-Applications and Q-components



© 2004 D. A. Menascé. All Rights Reserved.

58

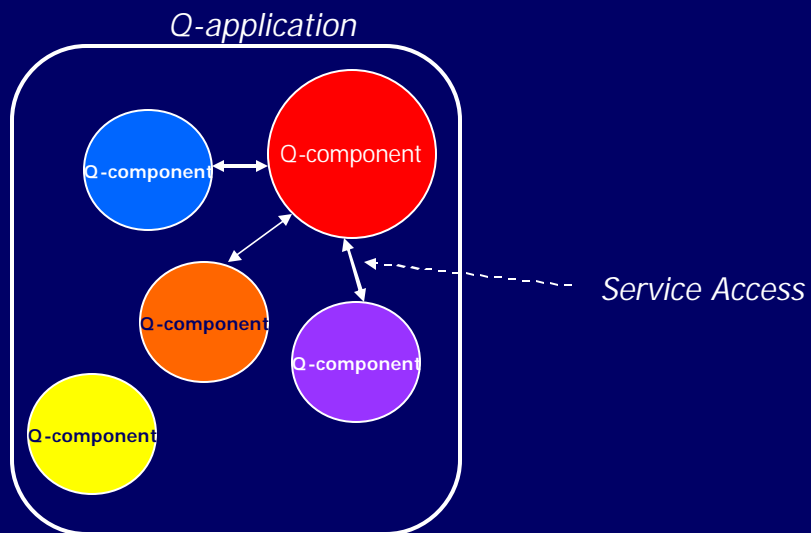
# Q-Applications and Q-components



© 2004 D. A. Menascé. All Rights Reserved.

59

# Q-Applications and Q-components



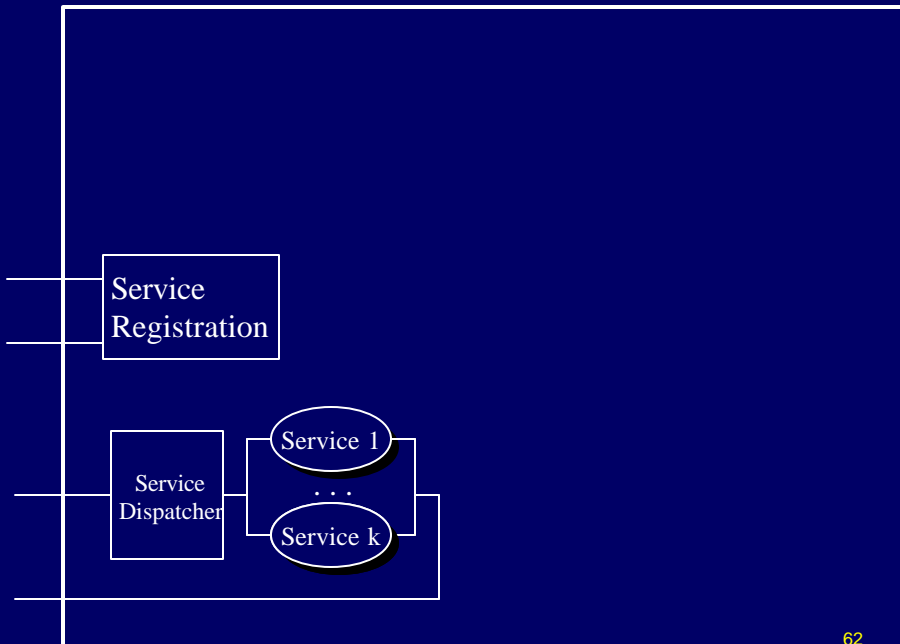
© 2004 D. A. Menascé. All Rights Reserved.

60

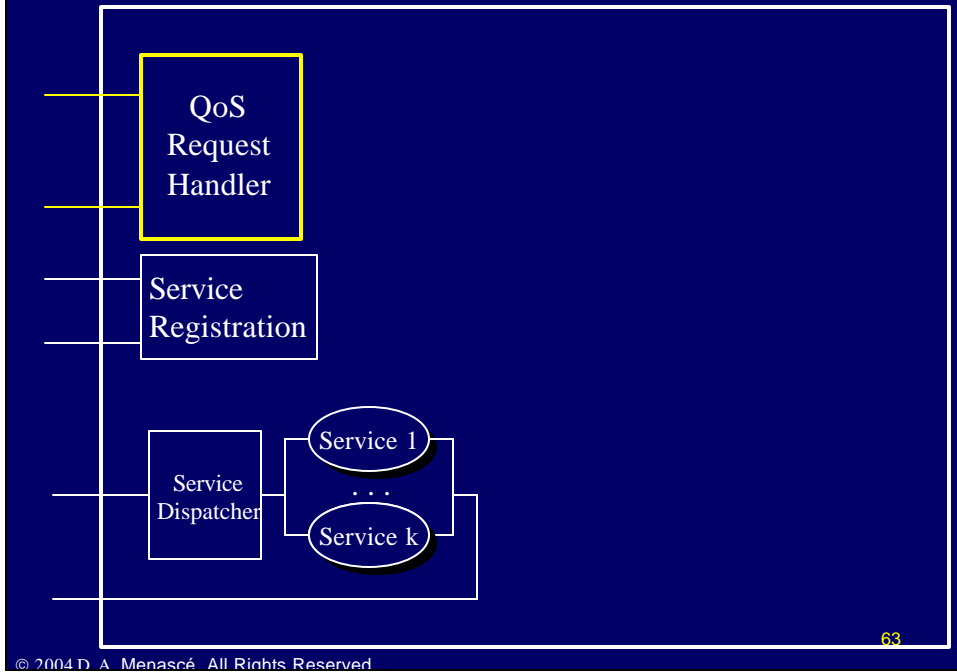
# QoS-Aware Software Components: Q-Components

- ❑ Engage in QoS Negotiations (accept, reject, counter-offer)
- ❑ Provide QoS guarantees for multiple concurrent services
- ❑ Maintain a table of QoS commitments
- ❑ Service dispatching based on accepted QoS commitments
- ❑ Q-components are the building blocks of QoS-aware applications

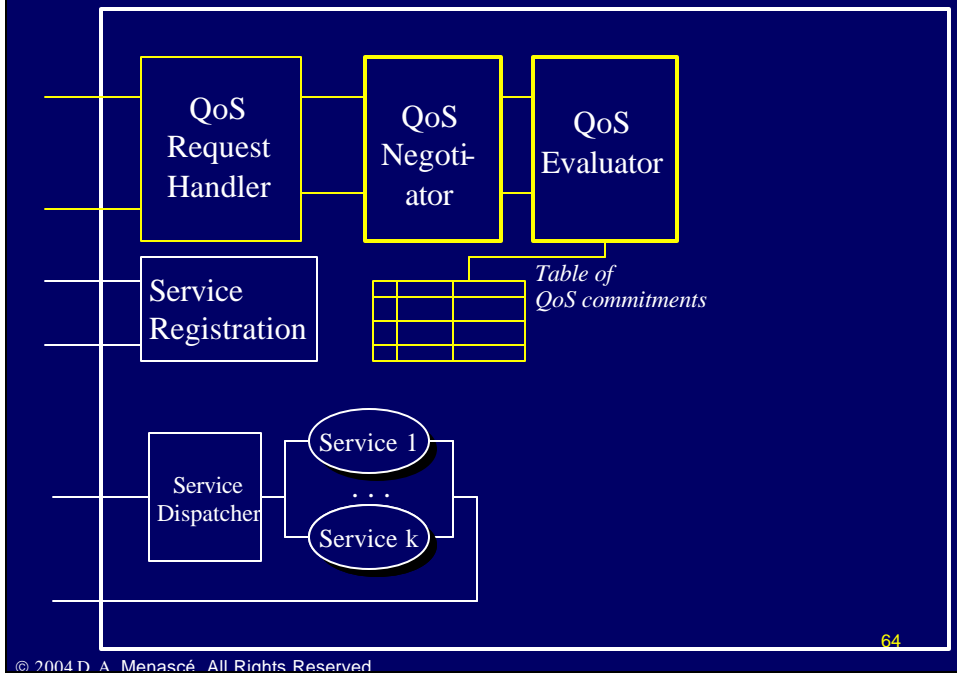
## *Architecture of a typical software component*



### Architecture of a Q-component (QoS Negotiation)

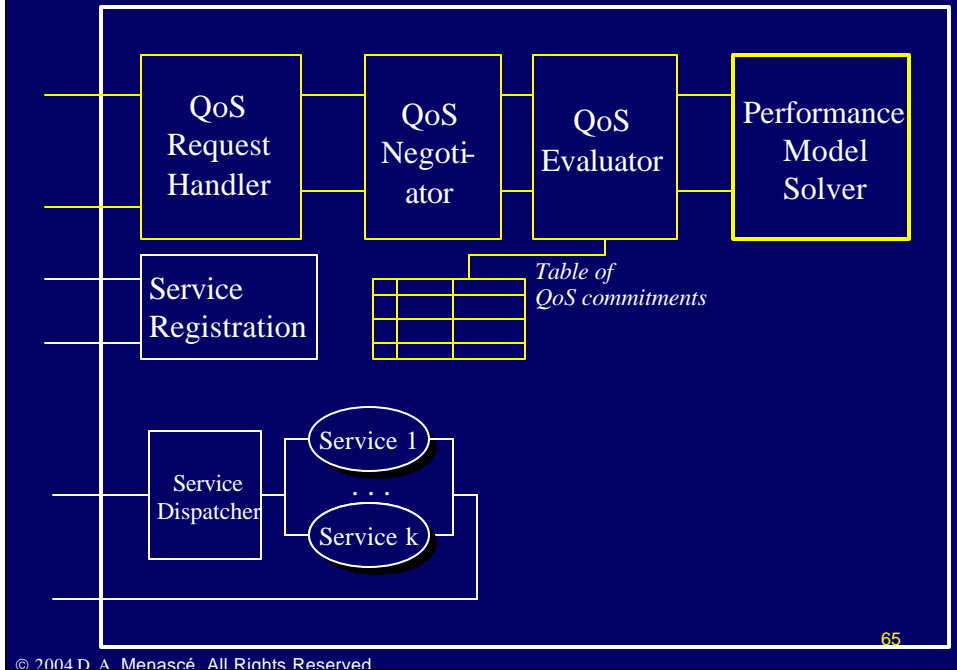


### Architecture of a Q-component (QoS Negotiation)



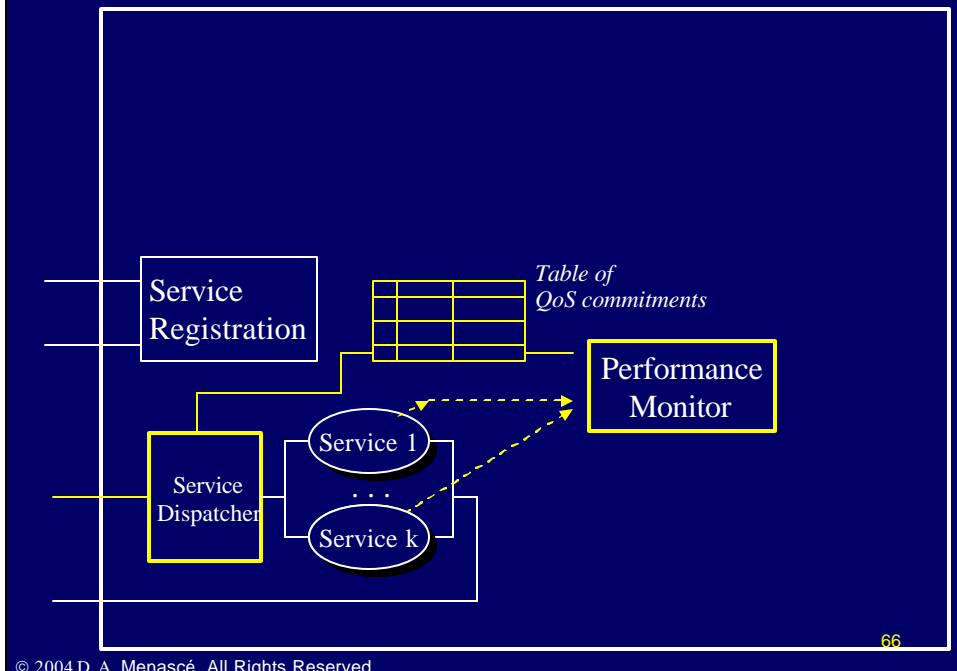


### Architecture of a Q-component (QoS Negotiation)



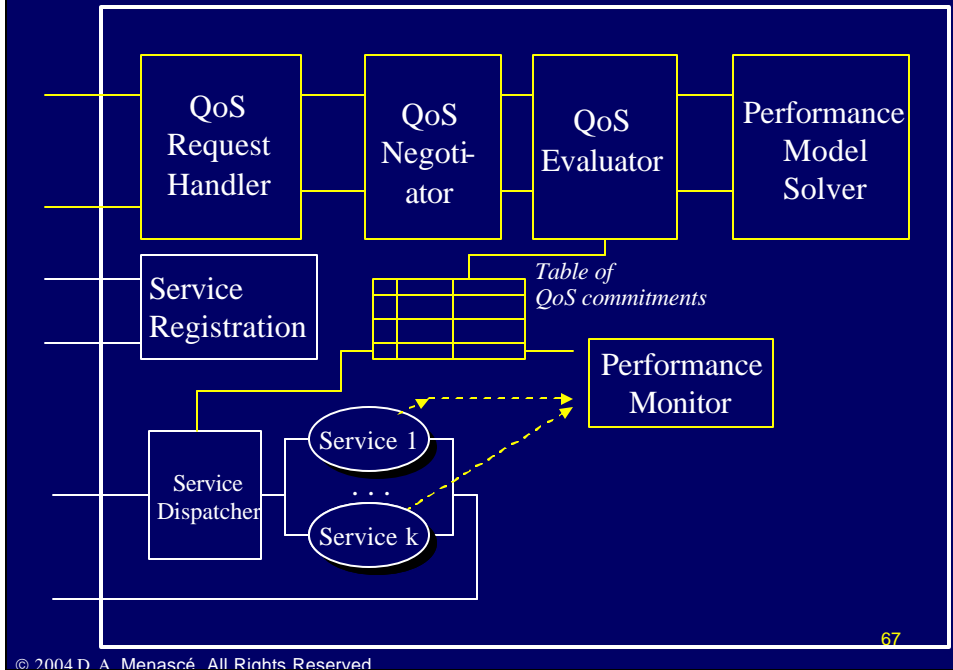
65

### Architecture of a Q-component –Service Requests



66

## Architecture of a Q-component



## Successful QoS Negotiation

Client

Q-component

QoSRequest(rid,Sid,N,Rmax,Xmin)

Accept(rid,token)

ServiceReq (... ,token)

ReplyReq (...)

ServiceReq (... ,token)

ReplyReq (...)

EndSession (token)

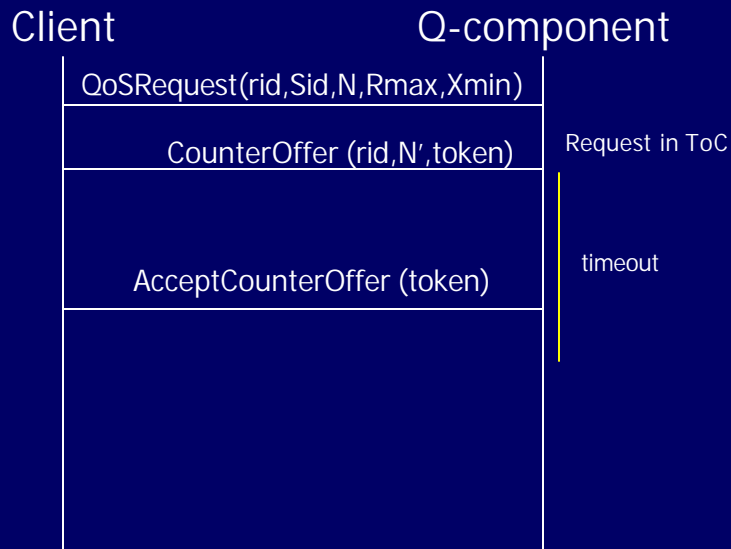
Request in ToC

Request removed from ToC

© 2004 D. A. Menascé. All Rights Reserved.

68

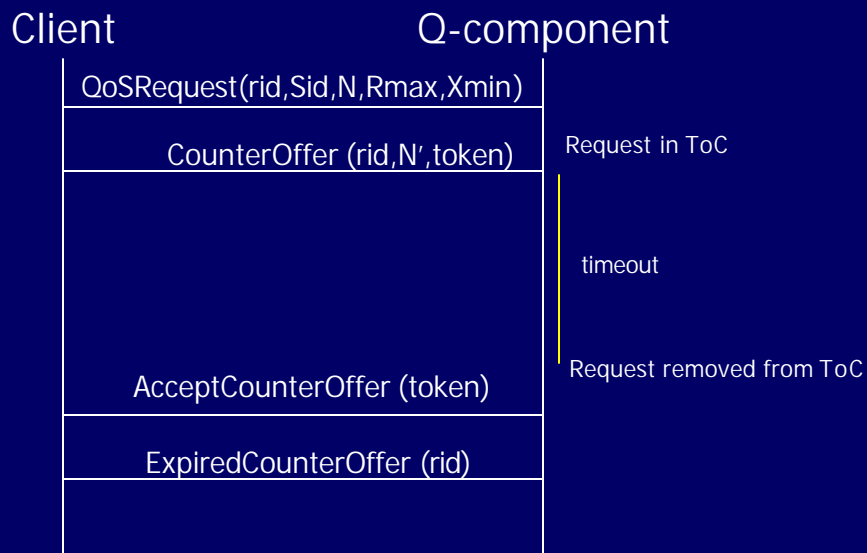
## On-time Accepted Counteroffer



© 2004 D. A. Menascé. All Rights Reserved.

69

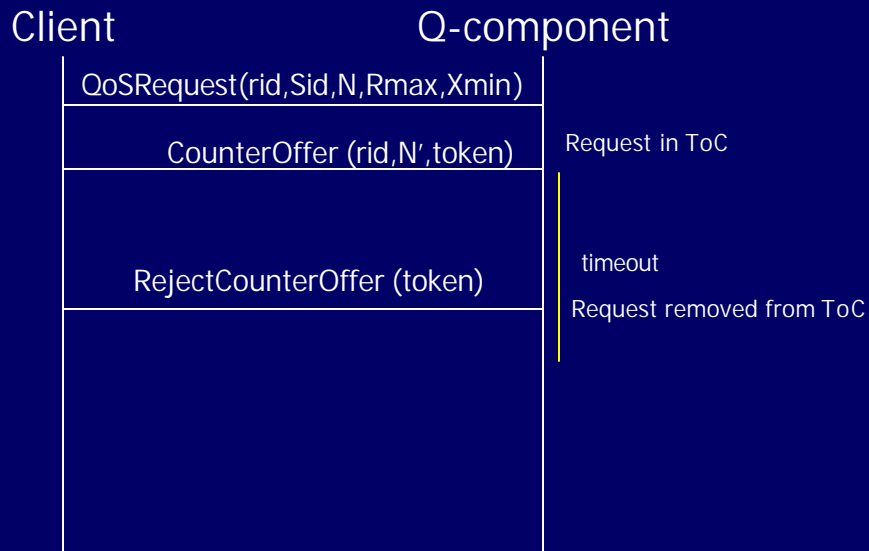
## Expired Accepted Counteroffer



© 2004 D. A. Menascé. All Rights Reserved.

70

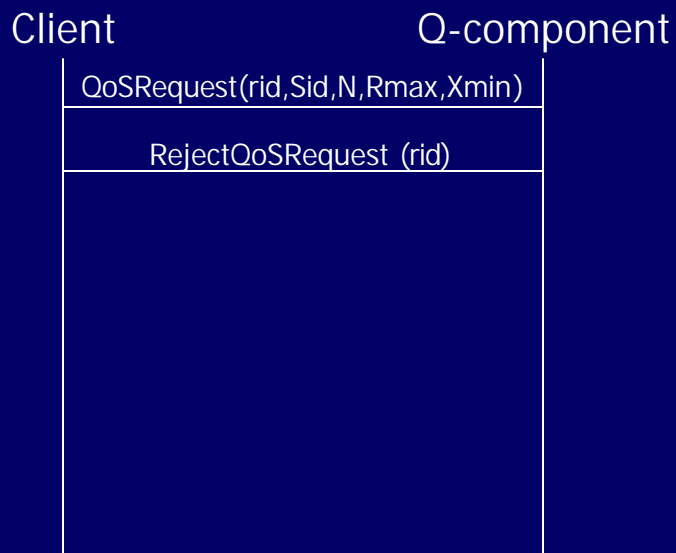
## Rejected Counteroffer



© 2004 D. A. Menascé. All Rights Reserved.

71

## Rejected QoS Negotiation



© 2004 D. A. Menascé. All Rights Reserved.


72

# Decision Table for QoS Negotiation

1. Current and other requests are satisfied				Accept
2. Only Current is Violated				
Reason	Remedy	Current	Others	Decision
Only MAXR is violated	Decrease N	OK	OK	Counter Offer
		Not OK: MINX is violated or N=0	OK	Reject
Only MINX is violated	Increase N	OK	OK	Counter Offer
		OK	Not OK	Reject
		Not OK & MAXR is violated	OK	Reject
		Not OK & MAXR is violated	Not OK	Reject
MINX & MAXR are violated	Decreasing N reduces X and increasing N increases R. So, there is no solution.			Reject
3. Current Request and Others are Violated				
Reason	Remedy	Current	Others	Decision
Only MAXR is violated	Decrease N	OK	OK	Counter Offer
		OK	Not OK	Reject
		Not OK: MINX is violated or N=0	OK or not OK	Reject
Only MINX is violated	X could be increased by increasing N. But this would further violate the QoS of other classes.			Reject
MINX & MAXR are violated	Decreasing N reduces X and increasing N increases R. So, there is no solution.			Reject
4. Only Other Requests are Violated				
Reason	Remedy	Current	Others	Decision
Any	Decrease N	OK	OK	Counter Offer
		OK	Not OK: N=1 but others still violated	Reject
		Not OK: MINX violated or N=0	OK or not OK	Reject

© 2004 D. A. Menascé. All Rights Reserved.

73



© 2004 D. A. Menascé. All Rights Reserved.

74

# Building a Performance Model

New Request: Sid = 3, N = 12

Base Matrix of Service Demands (in msec):

	Service		
	1	2	3
CPU	25	34	20
Disk 1	30	50	24
Disk 2	28	42	31

Table of Commitments (ToC):

Commitment ID	Service ID	N	...
1	2	10	...
2	3	15	...
3	1	8	...
4	1	20	...
5	2	13	...

Matrix of Service Demands (in msec)

	Class					
	1	2	3	4	5	6
CPU	34	20	25	25	34	20
Disk 1	50	24	30	30	50	24
Disk 2	42	31	28	28	42	31
Vector N:	10	15	8	20	13	12

© 2004 D. A. Menascé. All Rights Reserved.

75

# Building a Performance Model

New Request: Sid = 3, N = 12

Base Matrix of Service Demands (in msec):

	Service		
	1	2	3
CPU	25	34	20
Disk 1	30	50	24
Disk 2	28	42	31

Table of Commitments (ToC):

Commitment ID	Service ID	N	...
1	2	10	...
2	3	15	...
3	1	8	...
4	1	20	...
5	2	13	...


Matrix of Service Demands (in msec)

	Class					
	1	2	3	4	5	6
CPU	34	20	25	25	34	20
Disk 1	50	24	30	30	50	24
Disk 2	42	31	28	28	42	31
Vector N:	10	15	8	20	13	12

© 2004 D. A. Menascé. All Rights Reserved.

76


## Service 0 Results:



© 2004 D. A. Menascé. All Rights Reserved.

77


## Service 1 Results:



© 2004 D. A. Menascé. All Rights Reserved.

78

## Service 2 Results:



© 2004 D. A. Menascé. All Rights Reserved.

79


f = 0.0				
Svc No.	No. Dropped Sessions	No. of Sessions	% Drop	% Resp. Time Reduction
0	24	440	5	11
1	19	470	4	9
2	59	590	10	7
Total	102	1500	7	9
f = 0.10				
Svc No.	No. Dropped Sessions	No. of Sessions	% Drop	% Resp. Time Reduction
0	52	440	12	21
1	66	470	14	16
2	148	590	25	12
Total	266	1500	18	16
f = 0.25				
Svc No.	No. Dropped Sessions	No. of Sessions	% Drop	% Resp. Time Reduction
0	92	440	21	28
1	140	470	30	26
2	263	590	45	20
Total	495	1500	33	24

© 2004 D. A. Menascé. All Rights Reserved.

80




# Stock Quote Service



© 2004 D. A. Menascé. All Rights Reserved.

81

# Stock Quote Service



© 2004 D. A. Menascé. All Rights Reserved.

82

## Concluding Remarks

- ❑ Performance models can be used to build QoS controllers for complex multi-tiered systems:
  - Controlled system provides better QoS values even in case of high variability in request's inter-arrival and service times
  - Short term workload forecasting improves the QoS, especially when the workload intensity gets close to system saturation level
  - Dynamic adjustment of the controller interval length improves the QoS further
  - Even when basic model assumptions are violated, the models are robust enough to track the evolution of the performance metrics as the workload and configuration parameters change.

## Concluding Remarks (Cont'd)

- ❑ Performance models can be used by software components to make admission control decisions.
  - QoS components should be able to negotiate QoS requests and perform admission control
  - QoS negotiation overhead is small (it did not exceed 10% of the CPU service demand in our experiments).

# Bibliography

- “On the Use of Online Analytic Performance Models in Self-Managing and Self-Organizing Computer Systems,” D.A. Menascé, M. Bennanni and H. Ruan, in the book *Self-Star Properties in Complex Information Systems*, O. Babaoglu, M. Jelasity, A. Montresor, C. Fetzer, S. Leonardi, A. van Moorsel, and M. van Steen, eds., Lecture Notes in Computer Science, Vol. 3460, Springer Verlag, 2005.
- “Assessing the Robustness of Self-Managing Computer Systems under Highly Variable Workloads,” M. Bennani and D. Menascé, *Proc. International Conf. Autonomic Computing (ICAC-04)*, New York, NY, May 17-18, 2004.
- “A Framework for QoS-Aware Software Components,” D. Menascé, H. Ruan, and H. Gomaa, *Proc. 2004 ACM Workshop on Software and Performance (WOSP’04)*, San Francisco, CA, January 14, 2004.
- “On the Use of Performance Models to Design Self-Managing Systems,” D. Menascé and M. Bennani, *Proc. 2003 Computer Measurement Group Conference*, Dallas, TX, Dec. 71-2, 2003.
- “Automatic QoS Control,” *IEEE Internet Computing*, January/February 2003, Vol. 7, No. 1.
- “Preserving QoS of E-commerce Sites Through Self-Tuning: A Performance Model Approach,” D. A. Menascé, R. Dodge and D. Barbara, *Proc. 2001 ACM Conference on E-commerce*, Tampa, FL, October 14-17, 2001.